



Memory Management





Outline

- Concept of Memory
- Memory abstraction
- Logical and Physical address map
- Memory allocation
- Virtual Memory: Basics of Virtual Memory
- Paging
- Page Replacement Algorithms
- Segmentation


What is Memory?

- Computer memory is any **physical device capable of storing information temporarily or permanently**.
- Types of memory
 1. **Random Access Memory (RAM)**, is a **volatile memory** that **loses its contents when the computer** or hardware device **loses power**.
 2. **Read Only Memory (ROM)**, is a **non-volatile memory**, sometimes abbreviated as NVRAM, is a memory that **keeps its contents even if the power is lost**.
 - Computer **uses special ROM** called **BIOS (Basic Input Output System)** which **permanently stores the software** needed to access computer hardware such as hard disk and then load an operating system into RAM and start to execute it.



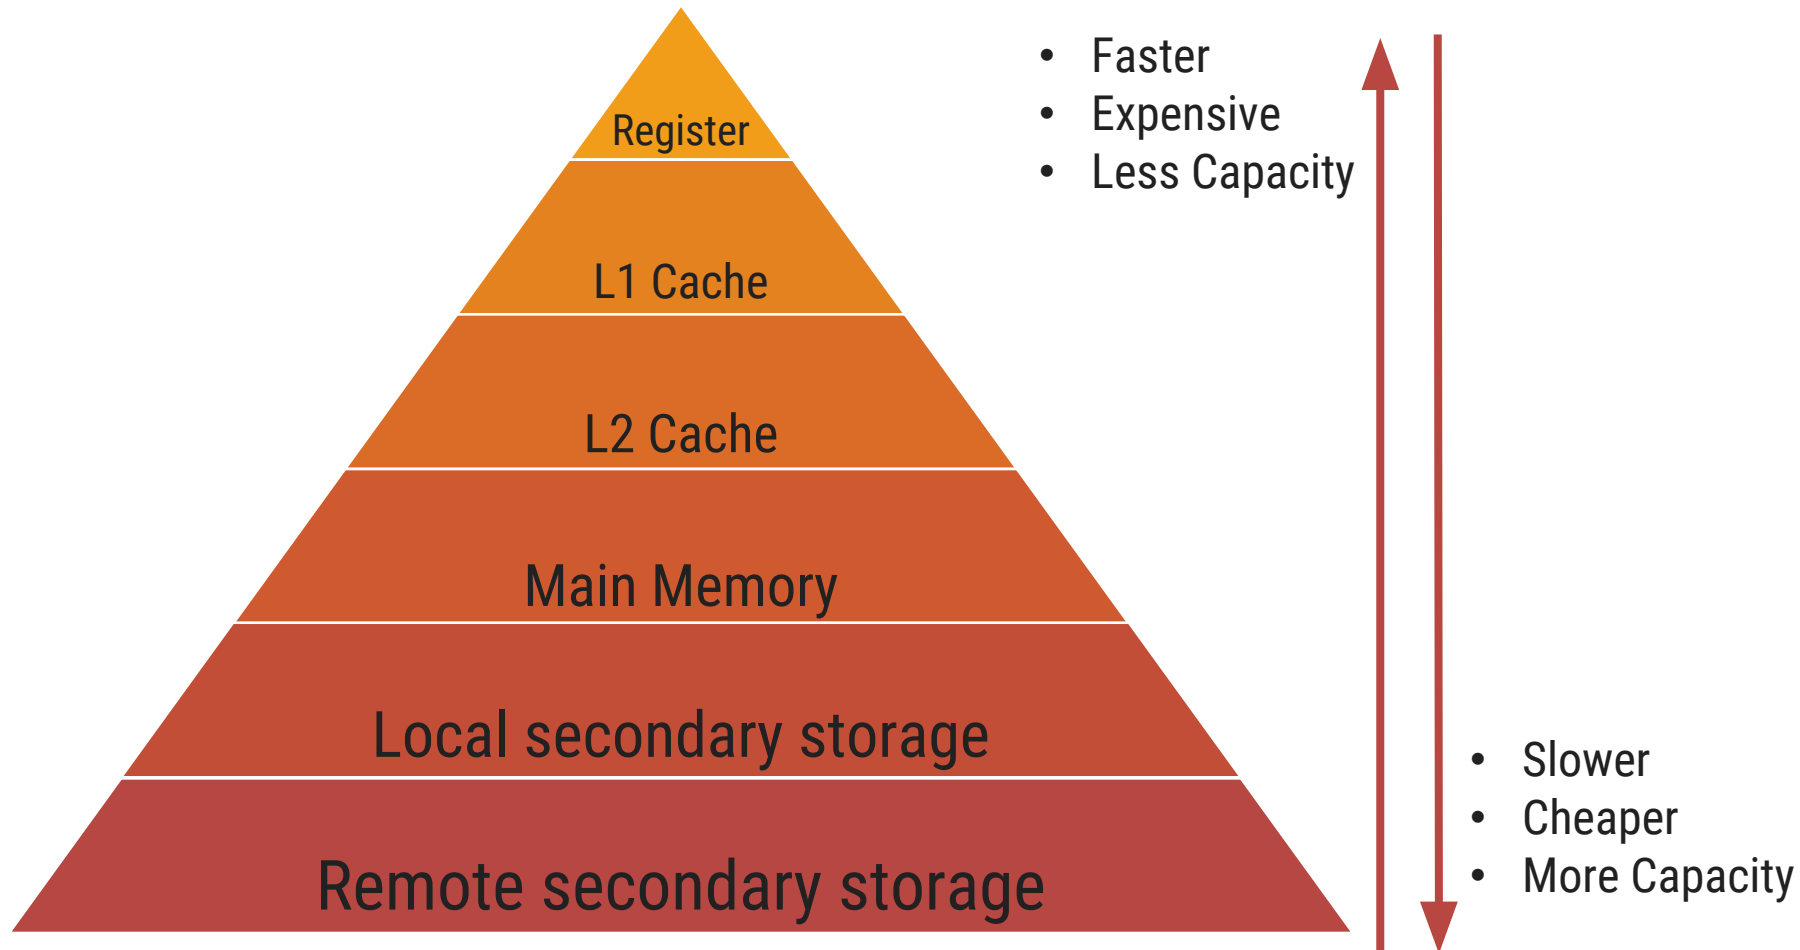
What is Memory?

- Computer memory is any **physical device capable of storing information temporarily or permanently**.
- Types of memory
 3. **Programmable Read-Only Memory (PROM)**, is a memory chip on which you can store a program. But once the PROM has been used, you **cannot wipe it clean and use it to store something else**. Like ROMs, PROMs are non-volatile. E.g CD-R
 4. **Erasable Programmable Read-Only Memory (EPROM)**, is a special type of PROM that **can be erased by exposing it to ultraviolet light**. E.g CD-RW
 5. **Electrically Erasable Programmable Read-Only Memory (EEPROM)**, is a special type of PROM that **can be erased by exposing it to an electrical charge**. E.g Pendrive



What is Memory Hierarchy?

- The **hierarchical arrangement of storage** in current computer architectures is called the memory hierarchy.

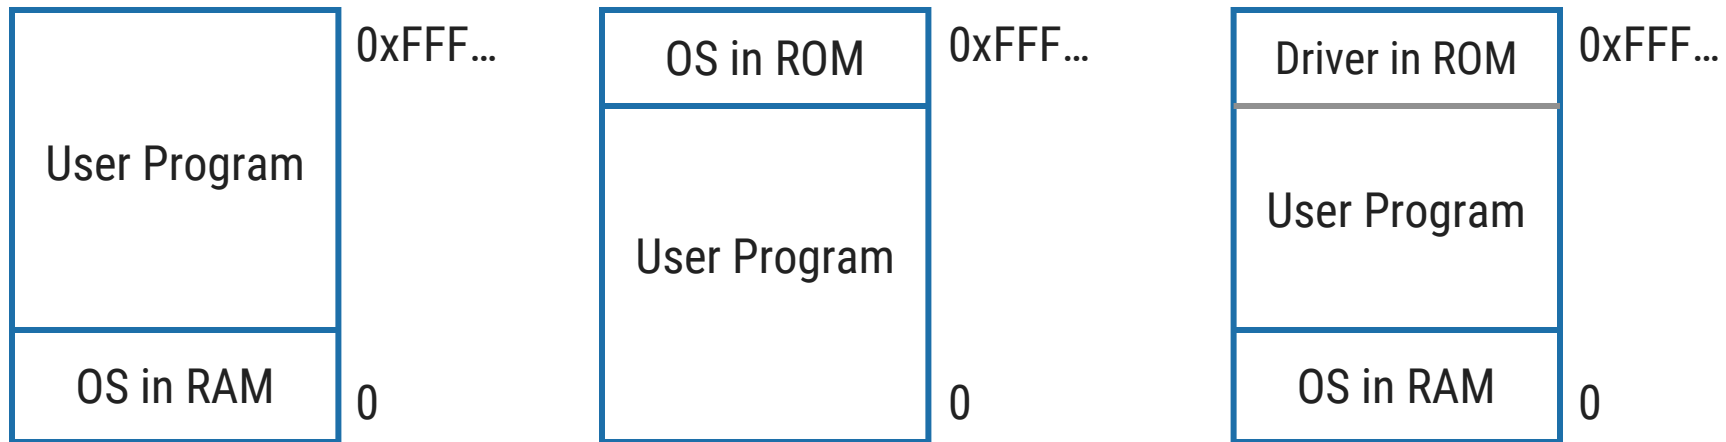


Memory abstraction

- The hardware and OS memory manager makes you see the memory as a single contiguous entity.
- How do they do that?
 - Abstraction
- Is abstraction necessary?

No memory abstraction


- In this model the memory presented to the programmer was simply a **single block of physical memory**.
 - having a set of addresses from 0 to some maximum
 - with each address corresponding to a cell containing some number of bits, commonly eight.
- When program execute instruction like **MOV REGISTER1, 1000**
- If at the **same time another program execute same instruction then value of first program will be overwrite**. So only **one process at a time can be running**.



Even with no abstraction, we can have several setups!

No memory abstraction

- What if we want to run multiple programs?
 1. OS saves entire memory on disk
 2. OS brings next program
 3. OS runs next program
- We can **use swapping** to **run multiple programs concurrently**.
- The **process of bringing in each process in its entirety in to memory, running it for a while and then putting it back on the disk** is called swapping.



Ways to implement swapping system

- Two different ways to implement swapping system
 1. Multiprogramming with **fixed partitions**
 2. Multiprogramming with **dynamic partitions**

Multiprogramming with fixed partitions

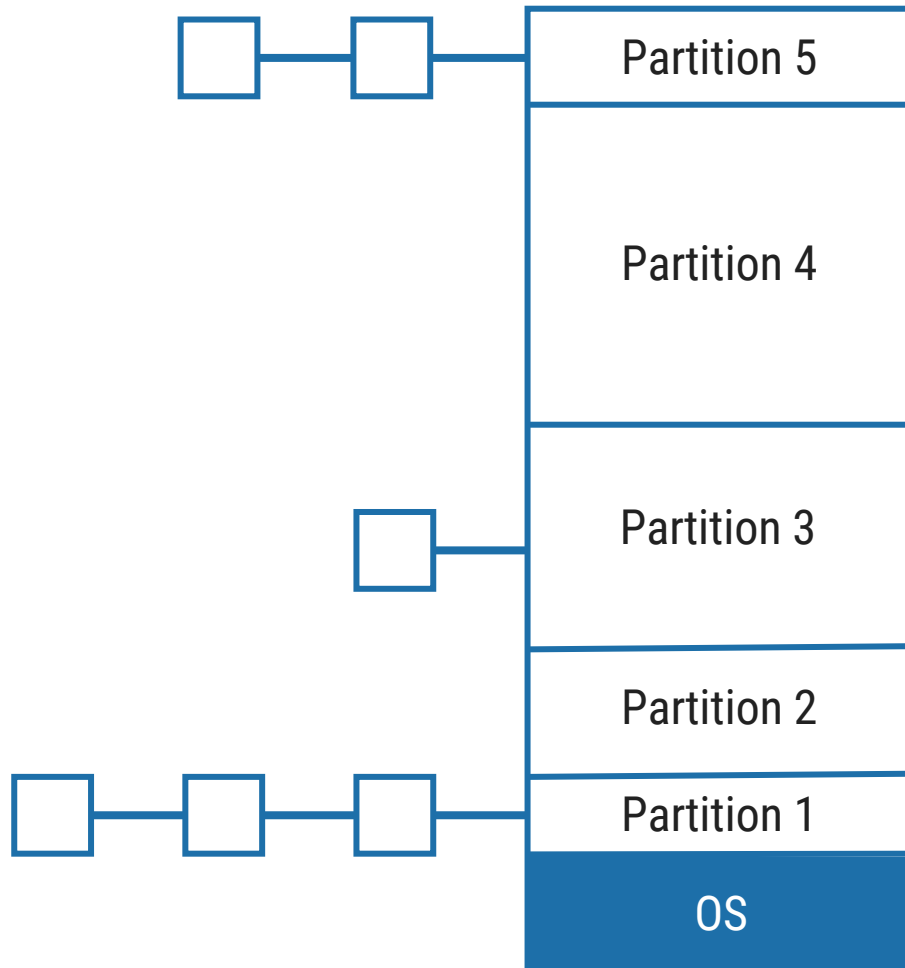
- ❑ Here memory is **divided into fixed sized partitions**.
- ❑ **Size can be equal or unequal** for different partitions.
- ❑ Generally **unequal partitions are used** for **better utilizations**.
- ❑ **Each partition can accommodate exactly one process**, means only single process can be placed in one partition.
- ❑ The partition **boundaries are not movable**.
- ❑ Whenever any **program needs to be loaded in memory**, a **free partition big enough to hold the program is found**. This **partition will be allocated** to that program or process.
- ❑ If **there is no free partition available of required size**, then the **process needs to wait**. Such **process will be put in a queue**.




Multiprogramming with fixed partitions

□ There are two ways to maintain queue

1. Using Multiple Input Queues

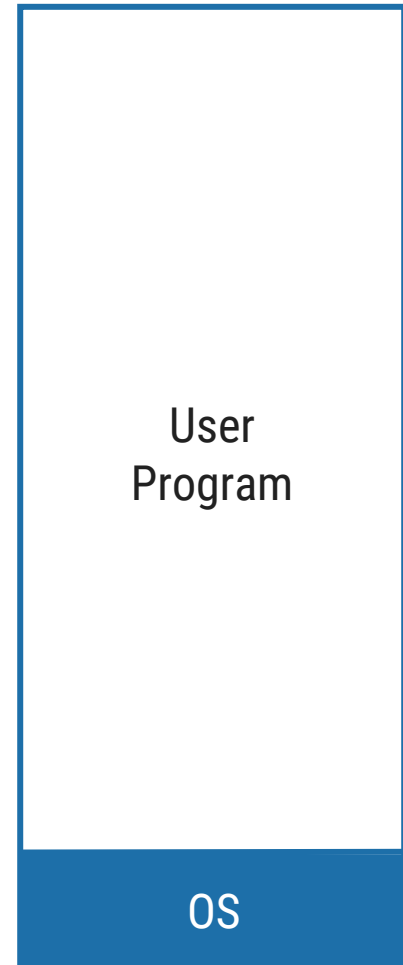


2. Using Single Input Queue

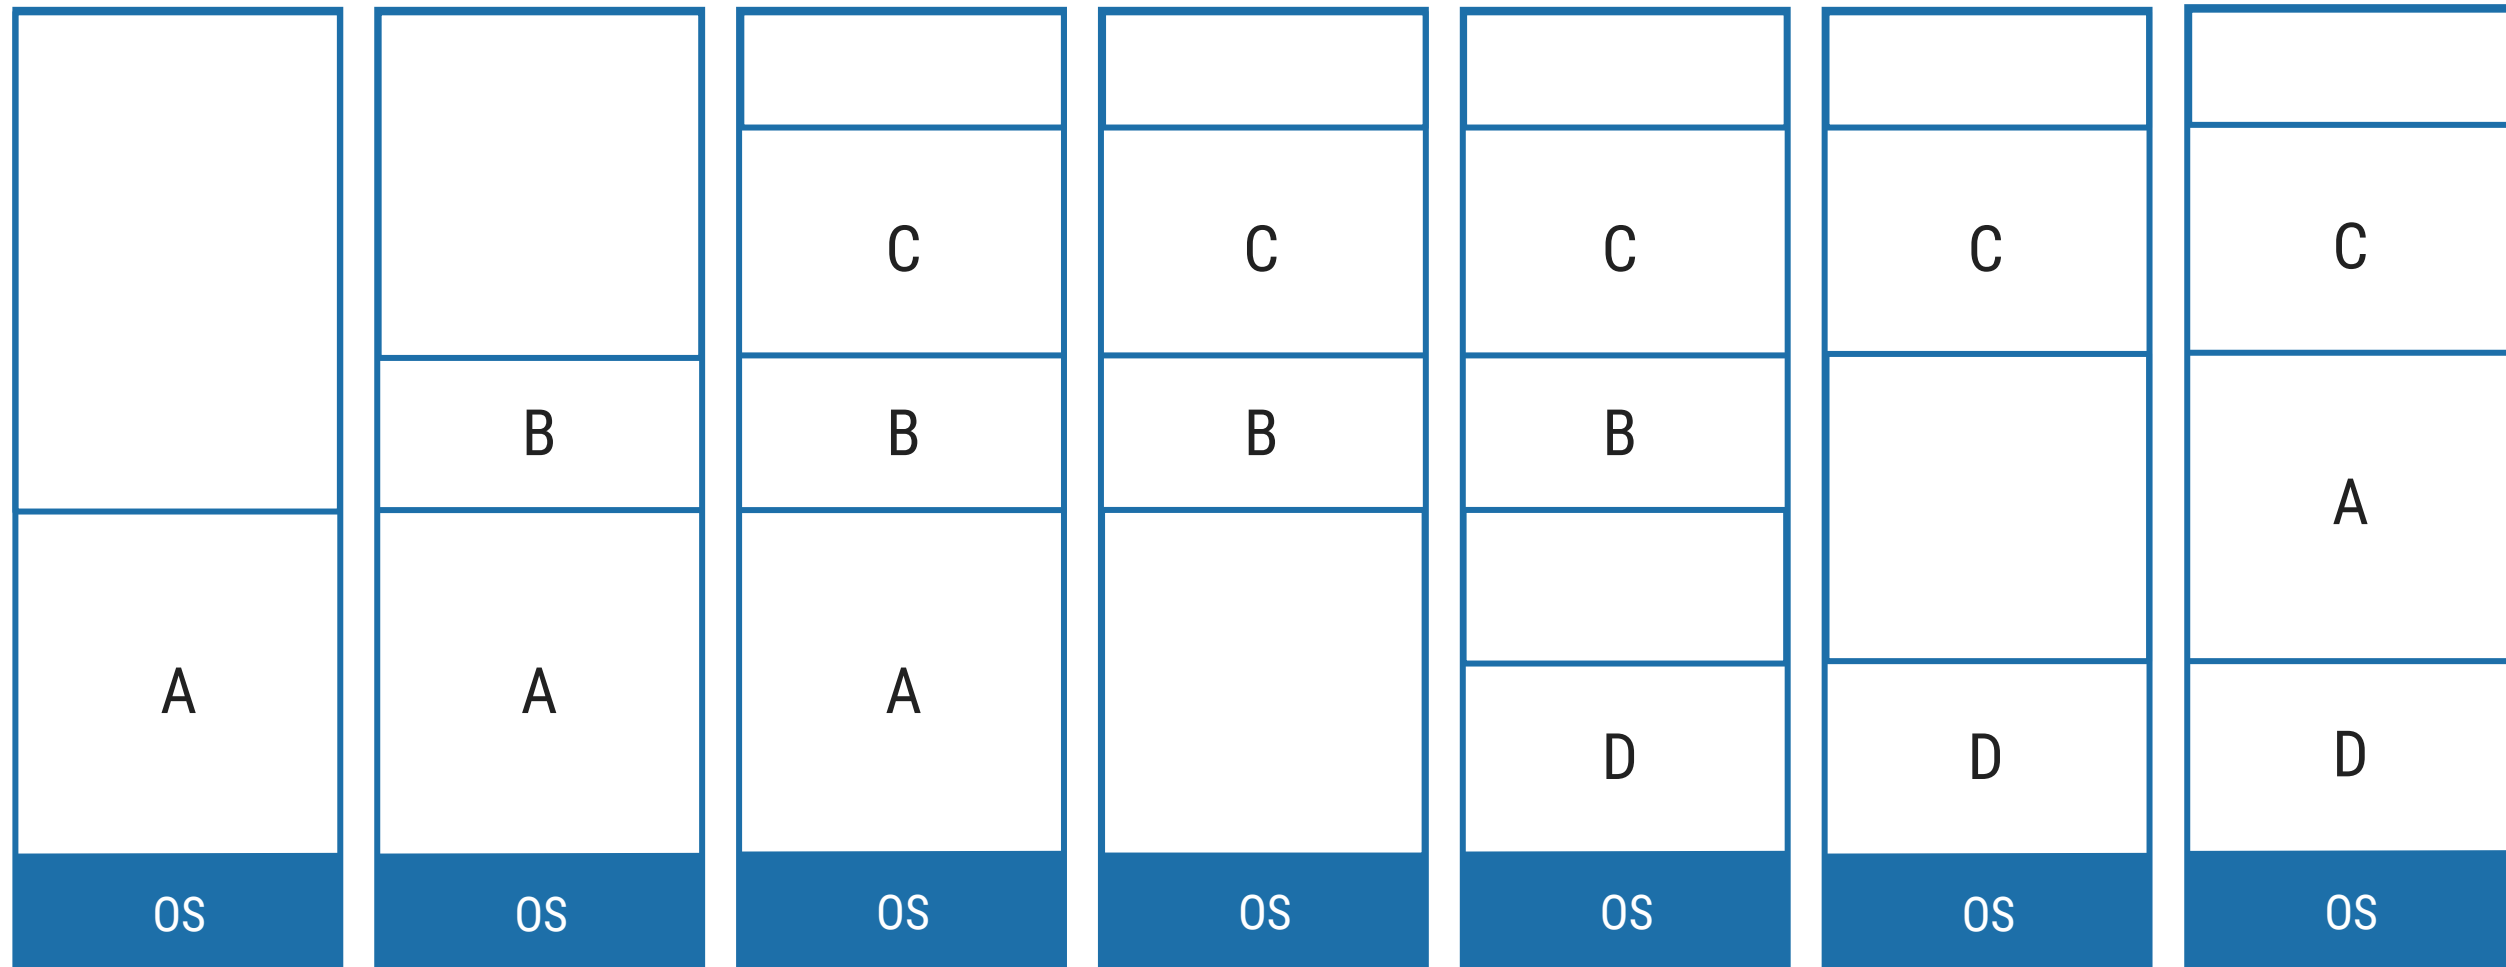


Multiprogramming with dynamic partitions

- ❑ Here, memory is **shared among operating system** and various simultaneously **running processes**.
- ❑ Initially, the **entire available memory is treated as a single free partition**.
- ❑ **Process is allocated exactly as much memory as it requires**.
- ❑ Whenever any **process enters** in a system, a **chunk of free memory big enough to fit the process is found and allocated**. The **remaining unoccupied space is treated as another free partition**.
- ❑ If **enough free memory is not available to fit the process, process needs to wait until required memory becomes available**.
- ❑ Whenever any **process gets terminate**, it **releases the space occupied**. If the **released free space is contiguous to another free partition, both the free partitions are merged together in to single free partition**.
- ❑ **Better utilization** of memory than fixed sized size partition.

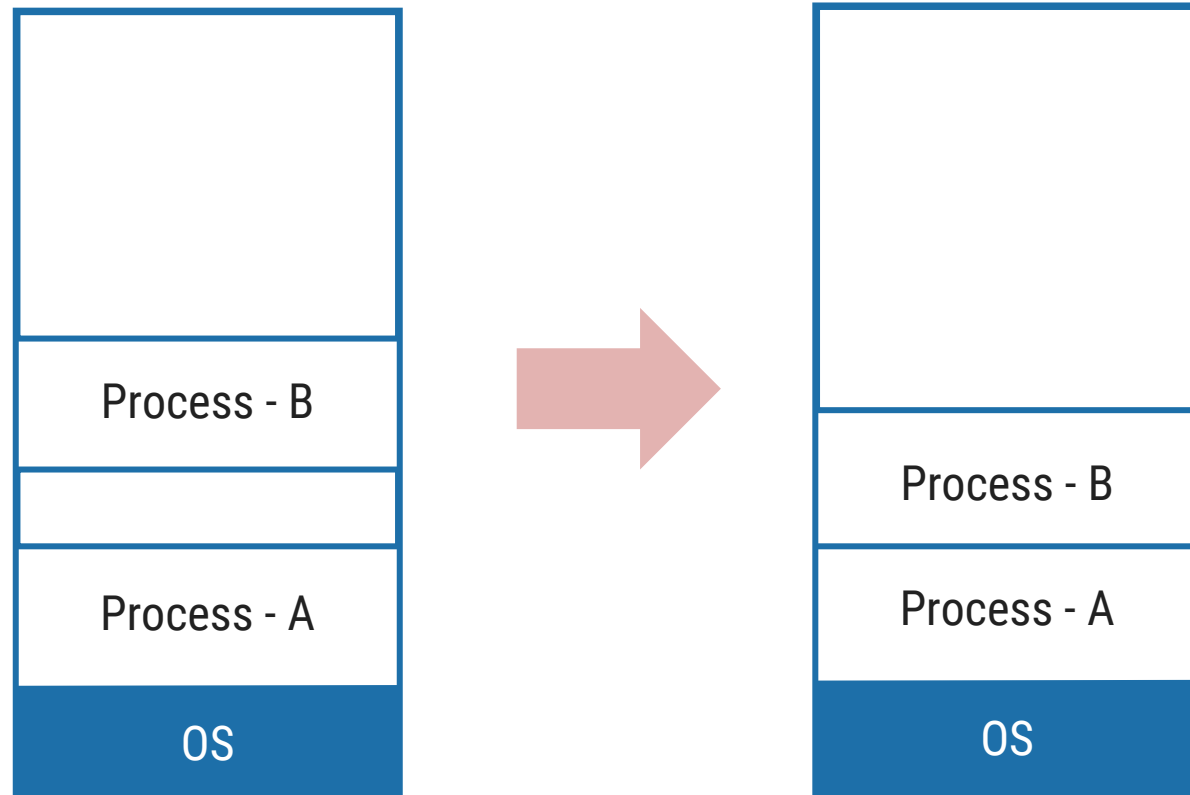


Multiprogramming with dynamic partitions

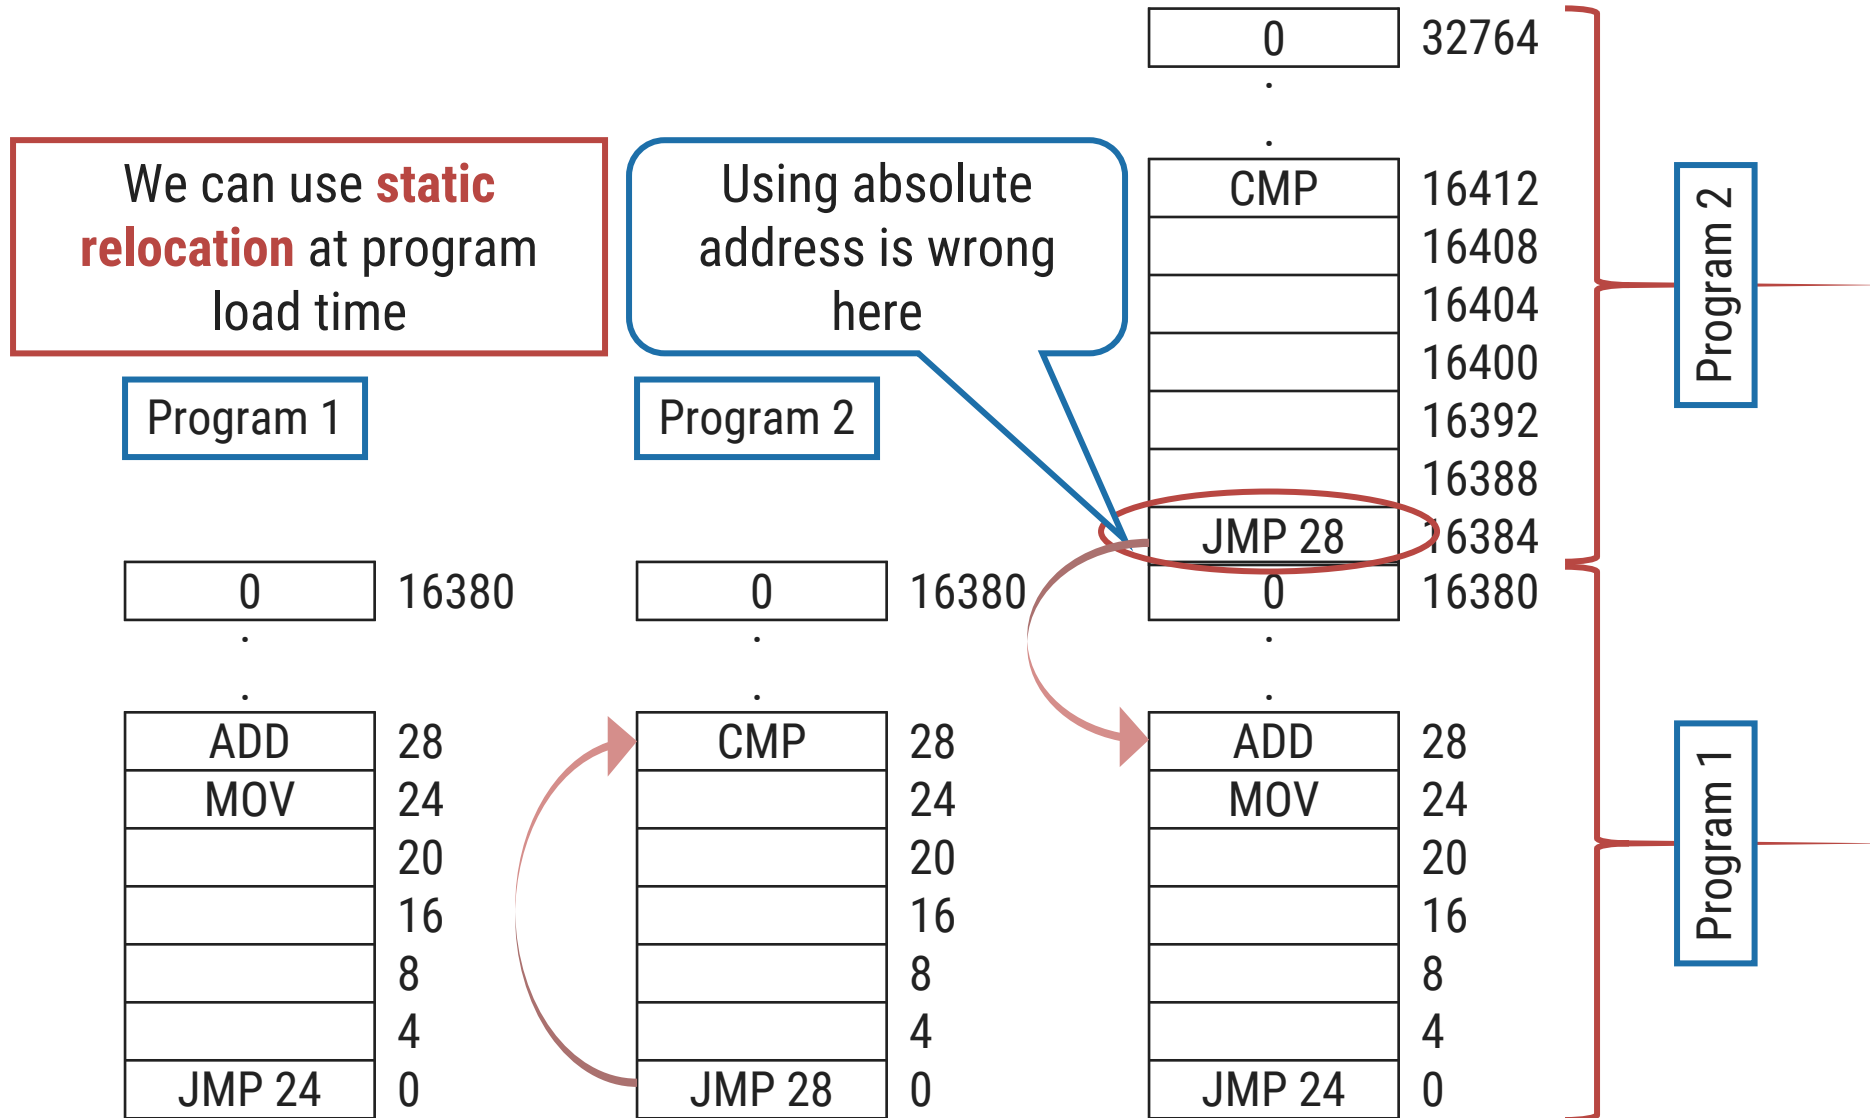


Memory compaction

- When **swapping creates multiple holes** in memory, it is possible to **combine them all in one big hole by moving all the processes downward** as far as possible. This technique is known as memory compaction.
- It **requires lot of CPU time**.



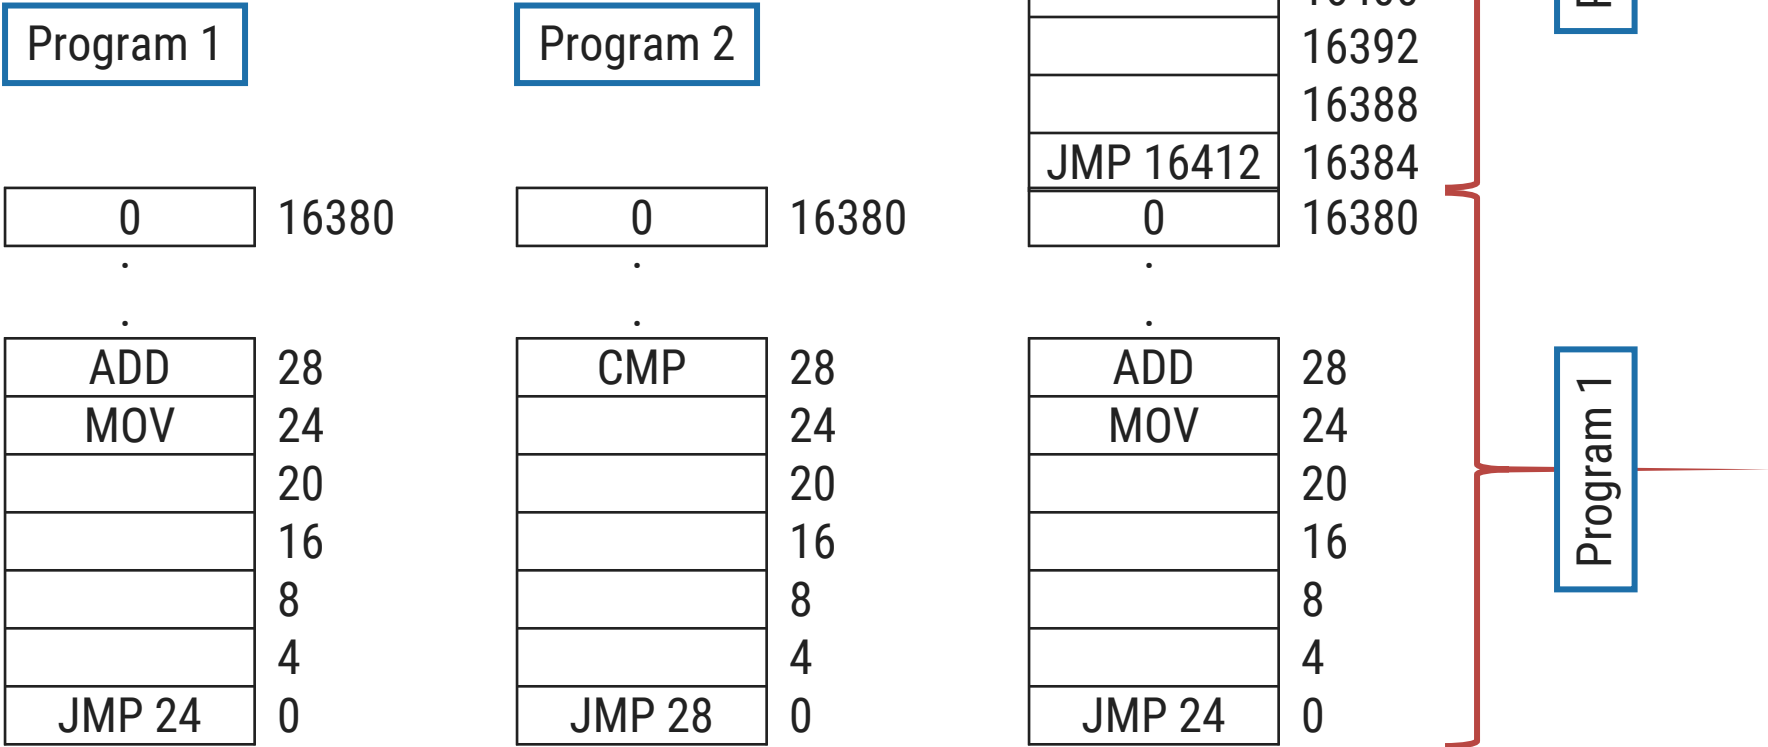
Multiprogramming without memory abstraction



Static relocation

When program was loaded at address 16384, the constant 16384 was added to every program address during the load process.

- Slow
- Required extra information from program



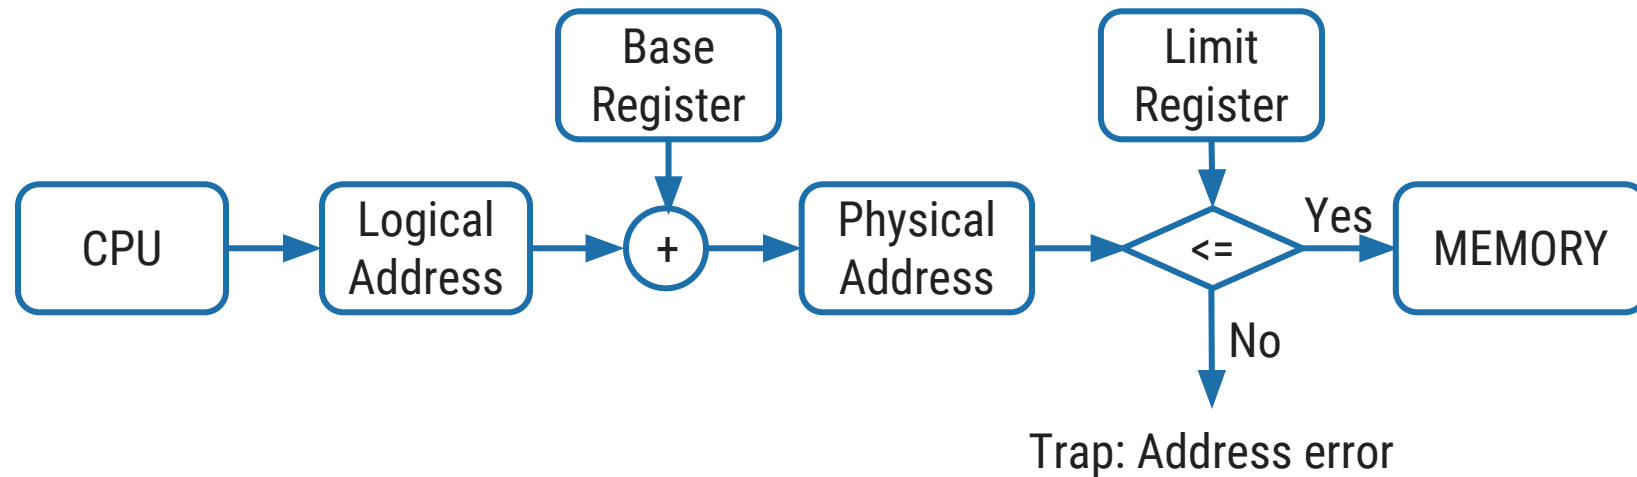
Base and Limit register

- ❑ An **address space** is set of addresses that a process can use to address **memory**.
- ❑ An **address space** is a range of valid addresses in memory that are **available for a program or process**.
- ❑ Two registers: **Base and Limit**
 - ❑ **Base register**: Start address of a program in physical memory.
 - ❑ **Limit register**: Length of the program.
- ❑ For every memory access
 - ❑ **Base is added to the address**
 - ❑ **Result compared to Limit**
- ❑ **Only OS can modify** Base and Limit register.

0	32764
.	.
CMP	16412
	16408
	16404
	16400
	16392
	16388
JMP 16412	16384
0	16380
.	.
ADD	28
MOV	24
	20
	16
	8
	4
JMP 24	0

Dynamic relocation

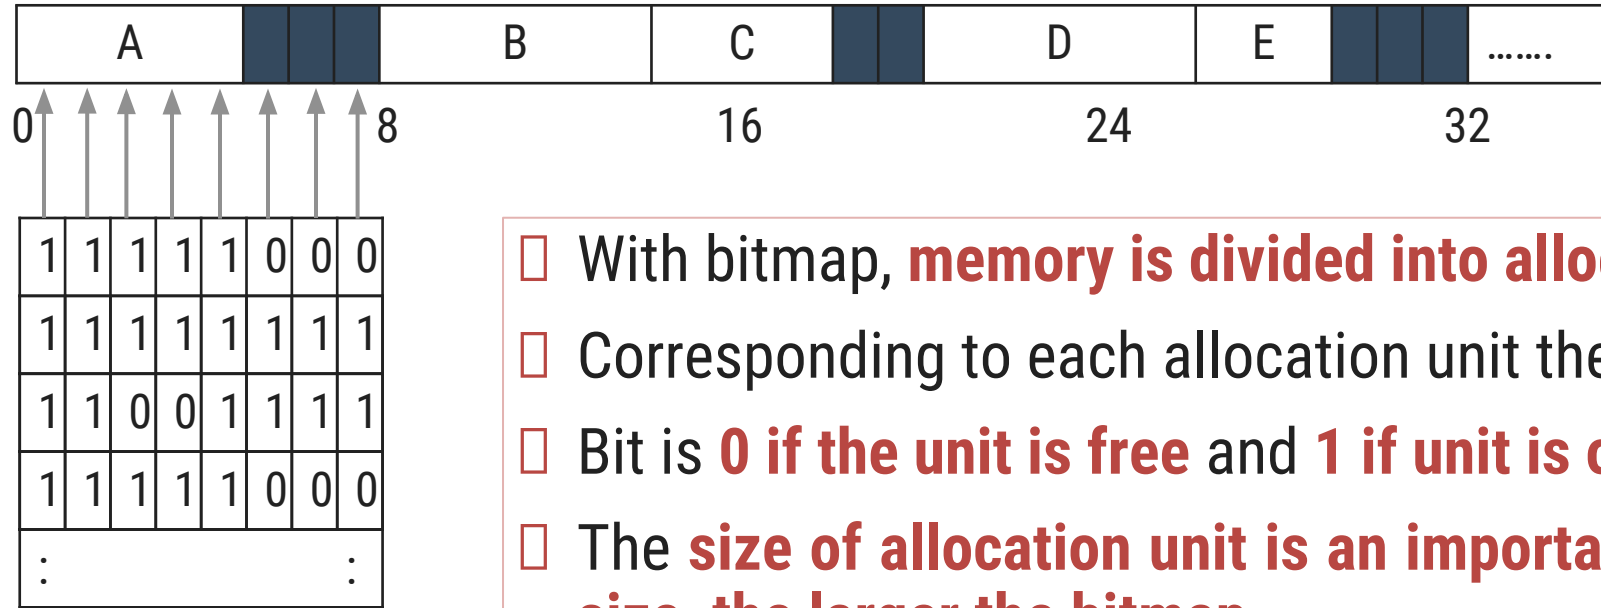
- Dynamic relocation refers to **address transformations being done during execution of a program.**
- Steps in dynamic relocation
 1. Hardware **adds relocation register (base)** to virtual address to get a physical address
 2. Hardware **compares address with limit register**; address **must be less than or equal limit**
 3. If **test fails**, the **processor takes an address trap** and **ignores the physical address.**



Managing free memory

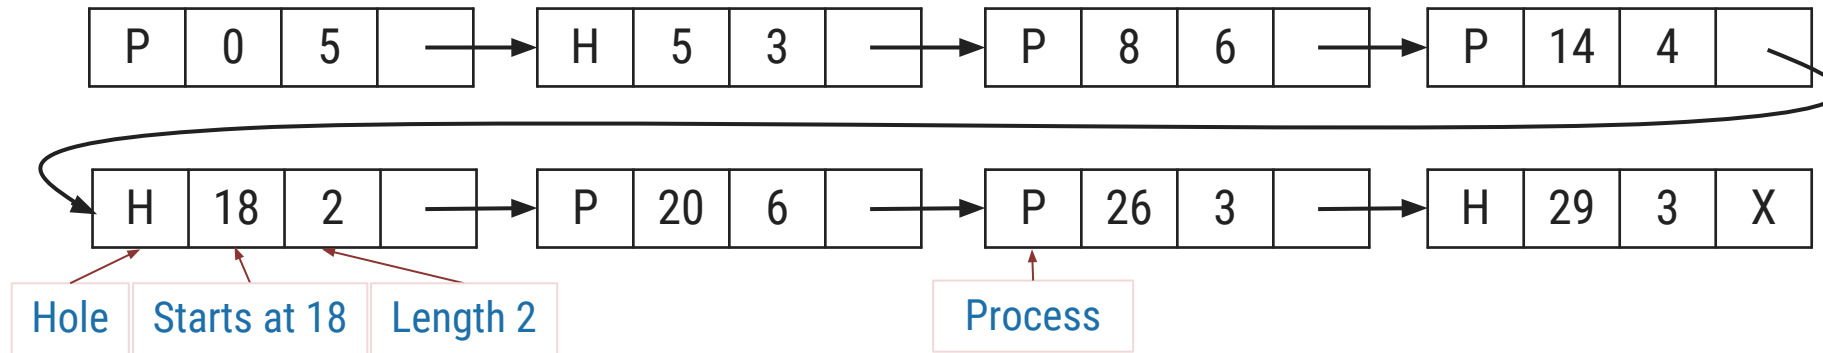
- Two ways to keep track of memory usage (free memory)
 1. Memory management with **Bitmaps**
 2. Memory management with **Linked List**

Memory management with Bitmaps




- ❑ With bitmap, **memory is divided into allocation units**.
- ❑ Corresponding to each allocation unit there is a **bit in a bitmap**.
- ❑ Bit is **0 if the unit is free** and **1 if unit is occupied**.
- ❑ The **size of allocation unit is an important design issue**, the **smaller the size, the larger the bitmap**.
- ❑ The main problem is that **when it has been decided to bring a k unit process, the memory manager must search the bitmap to find a run of k consecutive 0 bits in the map**.
- ❑ Searching a bitmap for a run of a given length is a **slow operation**.

Memory management with Linked List



- Another way to keep track of memory is **to maintain a linked list of allocated and free memory segments, where segment either contains a process or is an empty hole between two processes.**
- Each **entry in the list specifies a hole (H) or process (P), the address at which it starts the length and a pointer to the next entry.**
- The **segment list is kept sorted by address.**
- Sorting this way has the advantage that **when a process terminates or is swapped out, updating the list is straightforward.**
- A **terminating process normally has two neighbors** (except when it is at the very top or bottom of memory).

Memory management with Linked List



Before P terminate

After P terminate



P is replaced by H



P is replaced by H
and two H are merged



P is replaced by H
and two H are merged



P is replaced by H
and three H are merged



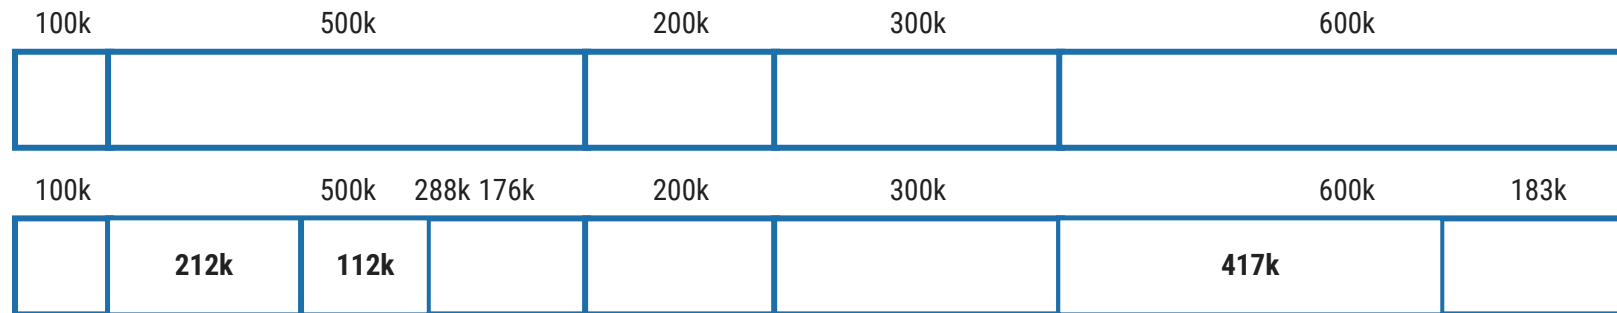
Memory allocation

□ Four memory allocation algorithms are as follow

1. First fit
2. Next fit
3. Best fit
4. Worst fit

First fit

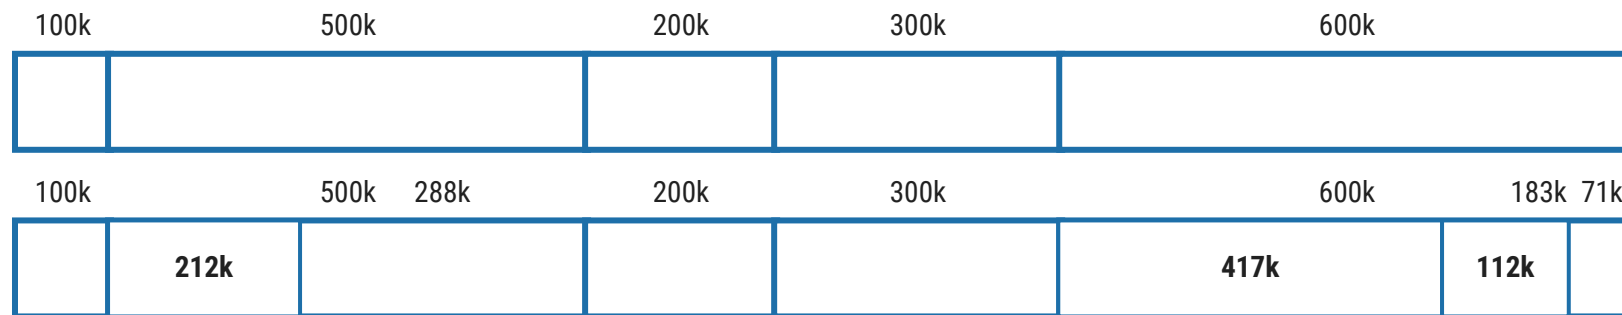
- ❑ Search **starts from the starting location** of the memory.
- ❑ **First available hole that is large enough to hold the process is selected** for allocation.
- ❑ The **hole is then broken up into two pieces**, one **for process** and another **for unused memory**.
- ❑ Example: Processes of **212K, 417K, 112K** and **426K** arrives in order.



- ❑ Here process of size **426k will not get any partition** for allocation.
- ❑ **Fastest algorithm** because it searches as little as possible.
- ❑ **Memory loss is higher**, as very large hole may be selected for small process.

Next fit


- It works in the same way as first fit, except that **it keeps the track of where it is whenever it finds a suitable hole.**
- The **next time when it is called to find a hole, it starts searching the list from the place where it left off last time.**
- Example: Processes of **212K, 417K, 112K** and **426K** arrives in order.



- Here process of size **426k will not get any partition** for allocation.
- **Search time is smaller.**
- Memory manager must have to **keep track of last allotted hole to process.**
- It **gives slightly worse performance** than first fit.

Best fit

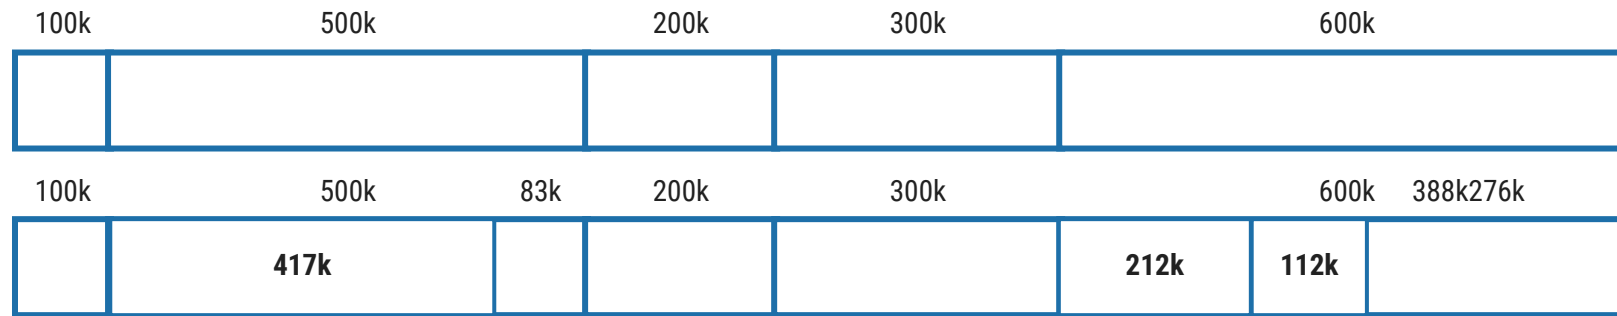
- ❑ Entire memory is searched here.
- ❑ The **smallest hole, which is large enough to hold the process, is selected** for allocation.
- ❑ Example: Processes of **212K, 417K, 112K** and **426K** arrives in order.



- ❑ **Search time is high**, as it searches entire memory every time.
- ❑ **Memory loss is less.**

Worst fit

- ❑ Entire memory is searched here also. The **largest hole, which is largest enough to hold the process, is selected** for allocation.
- ❑ Example: Processes of **212K, 417K, 112K** and **426K** arrives in order.




- ❑ Here process of size **426k will not get any partition** for allocation.
- ❑ **Search time is high**, as it searches entire memory every time.
- ❑ This algorithm can be **used only with dynamic partitioning**.

Virtual Memory

- ❑ Memory is **hardware that your computer uses to load the operating system and run programs.**
- ❑ Computer **consists of one or more RAM chips** that each have several memory modules.
- ❑ The amount of real **memory in a computer is limited** to the amount of RAM installed. Common memory sizes are **1GB, 2GB, and 4GB.**
- ❑ Because your computer has a finite amount of RAM, **it is possible to run out of memory when too many programs are running at one time.**
- ❑ This is where **virtual memory comes into the picture.**
- ❑ Virtual memory **increases the available memory of your computer by enlarging the "address space," or places in memory where data can be stored.**
- ❑ It does this by **using hard disk space** for additional memory allocation.
- ❑ However, since the **hard drive is much slower than the RAM, data stored in virtual memory must be mapped back to real memory in order to be used.**

Virtual Memory

- ❑ Each program has its own address space, which is **broken up into pages**.
- ❑ Each **page** is a contiguous range of addresses.
- ❑ These **pages** are mapped onto the **physical memory** but, **to run the program, all pages are not required** to be present in the physical memory.
- ❑ The operating system **keeps those parts of the program currently in use in main memory**, and the rest on the disk.
- ❑ In a system using virtual memory, the **physical memory is divided into page frames** and the **virtual address space is divided into equally-sized partitions called pages**.
- ❑ Virtual memory **works fine in a multiprogramming system**, with bits and pieces of many programs in memory at once.



Paging

- Paging is a **storage mechanism used to retrieve processes from the secondary storage (Hard disk) into the main memory (RAM) in the form of pages.**

Virtual
Address
Space

40K – 44K	Page 11
36K – 40K	Page 10
32K – 36K	Page 9
28K – 32K	Page 8
24K – 28K	Page 7
20K – 24K	Page 6
16K – 20K	Page 5
12K – 16K	Page 4
8K – 12K	Page 3
4K – 8K	Page 2
0K – 4K	Page 1

Virtual page

- The main idea behind the paging is **to divide each process in the form of pages**. The **main memory will also be divided in the form of frames**.
- One page of the process is to be stored in one of the frames** of the memory.

Page frame

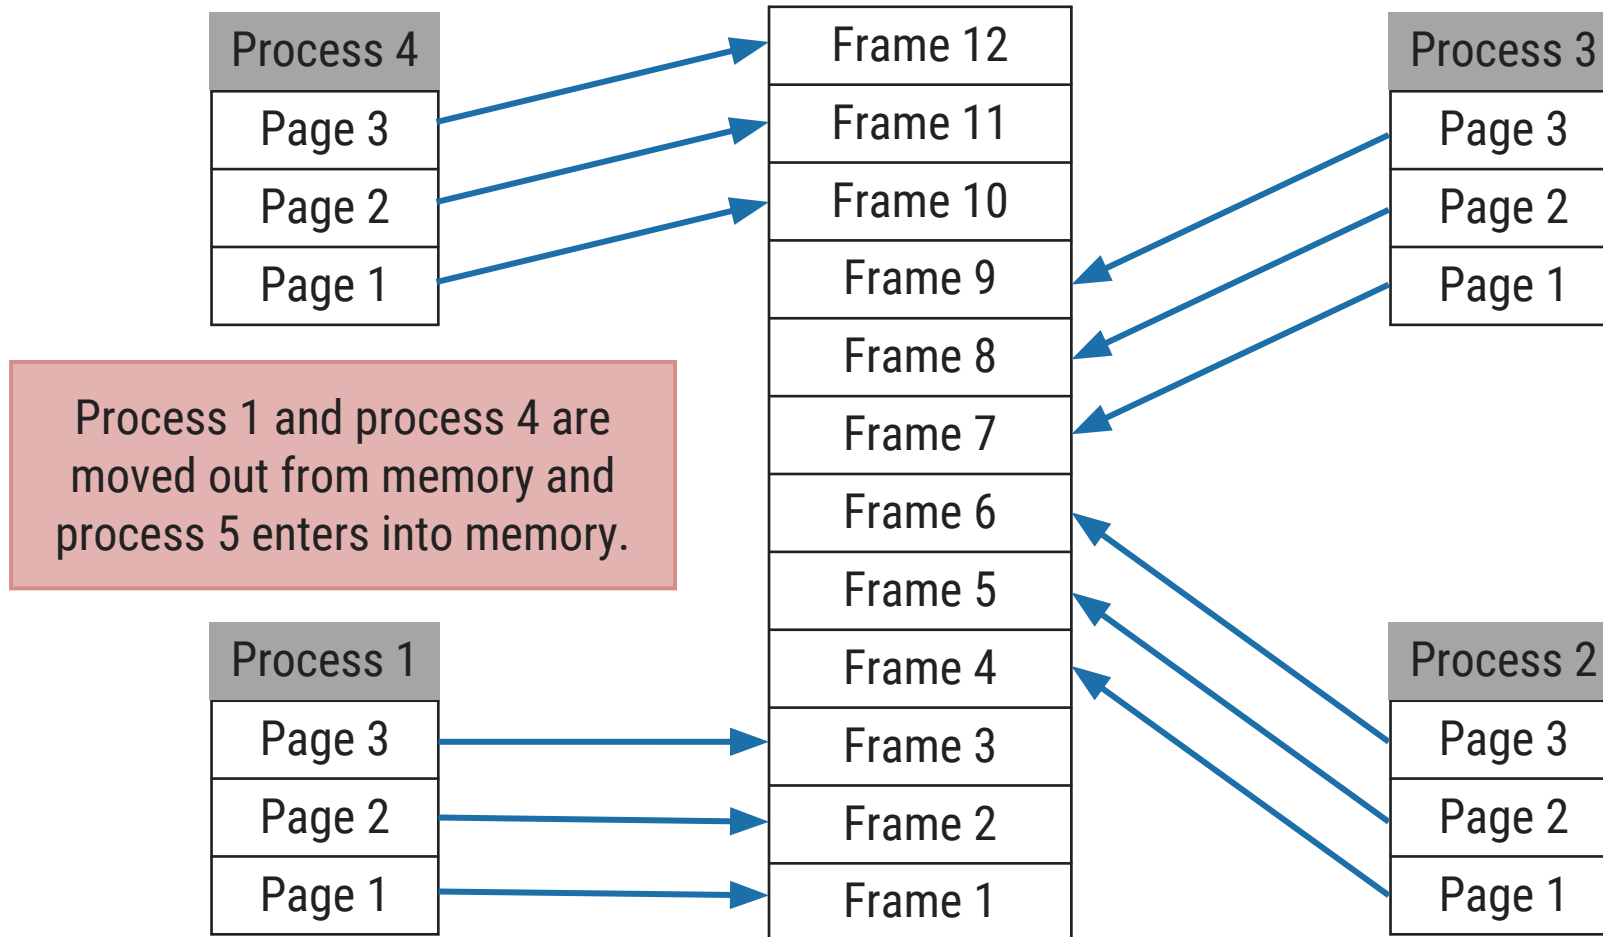
Frame 6	20K – 24K
Frame 5	16K – 20K
Frame 4	12K – 16K
Frame 3	8K – 12K
Frame 2	4K – 8K
Frame 1	0K – 4K

Physical
Memory
Address

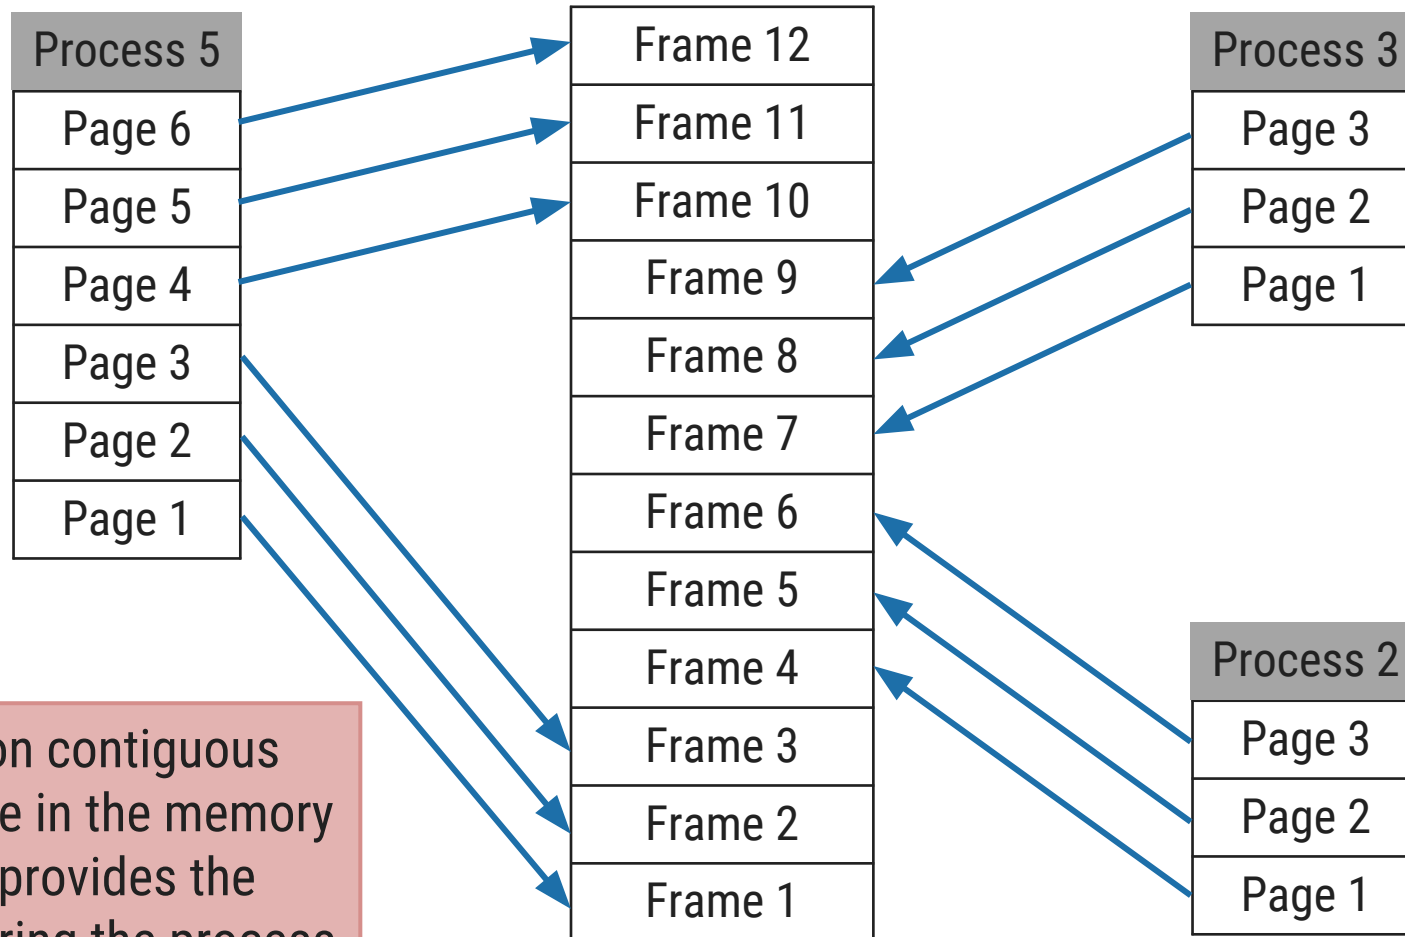
Paging

- ❑ The pages can be stored at the different locations of the memory but the priority is always to find the contiguous frames or holes.
- ❑ Pages of the process are **brought into the main memory only when they are required** otherwise they reside in the secondary storage.
- ❑ The **sizes of each frame must be equal**. Considering the fact that the pages are mapped to the frames in Paging, page size needs to be as same as frame size.
- ❑ Different operating system defines different frame sizes.

Paging



Paging



We have 6 non contiguous frames available in the memory and paging provides the flexibility of storing the process at the different places.

Paging

- **Size of Virtual Address Space is greater than that of main memory**, so instead of loading entire address space in to memory to run the process, **MMU copies only required pages into main memory**.
- In order to **keep the track of pages and page frames**, OS maintains a data structure called **page table**.

Virtual
Address
Space

40K – 44K	
36K – 40K	
32K – 36K	5
28K – 32K	
24K – 28K	
20K – 24K	3
16K – 20K	4
12K – 16K	0
8K – 12K	
4K – 8K	1
0K – 4K	2

Virtual page

Page frame

5	20K – 24K
4	16K – 20K
3	12K – 16K
2	8K – 12K
1	4K – 8K
0	0K – 4K

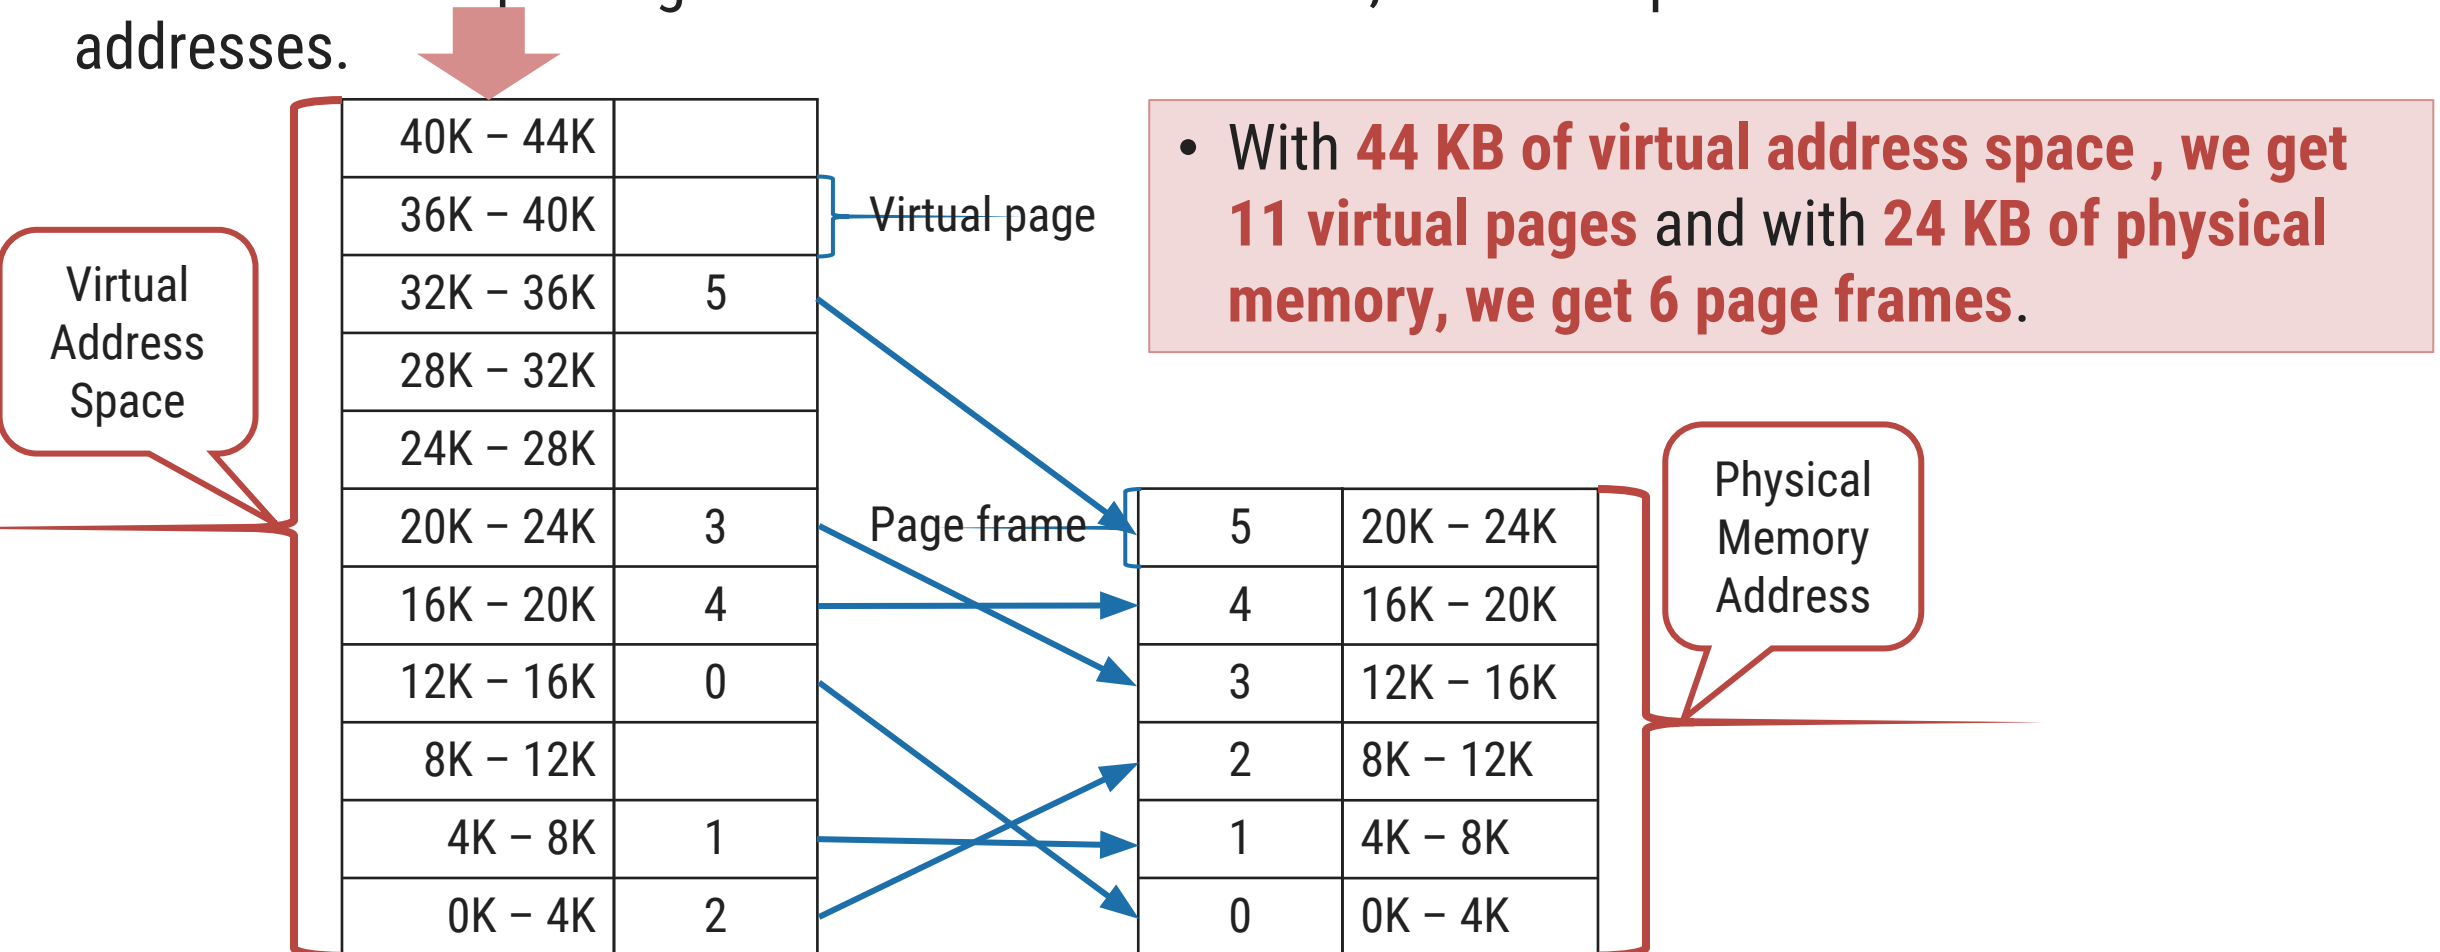
Physical
Memory
Address

Logical Address vs Physical Address

- ❑ **Logical Address is generated by CPU** while a program is running. The **logical address is virtual address as it does not exist physically** therefore it is also known as **Virtual Address**.
- ❑ **Physical Address identifies a physical location of required data in a memory**. The user never directly deals with the physical address but can access by its corresponding logical address.
- ❑ The **hardware device called Memory-Management Unit is used for mapping (converting) logical address to its corresponding physical address**.

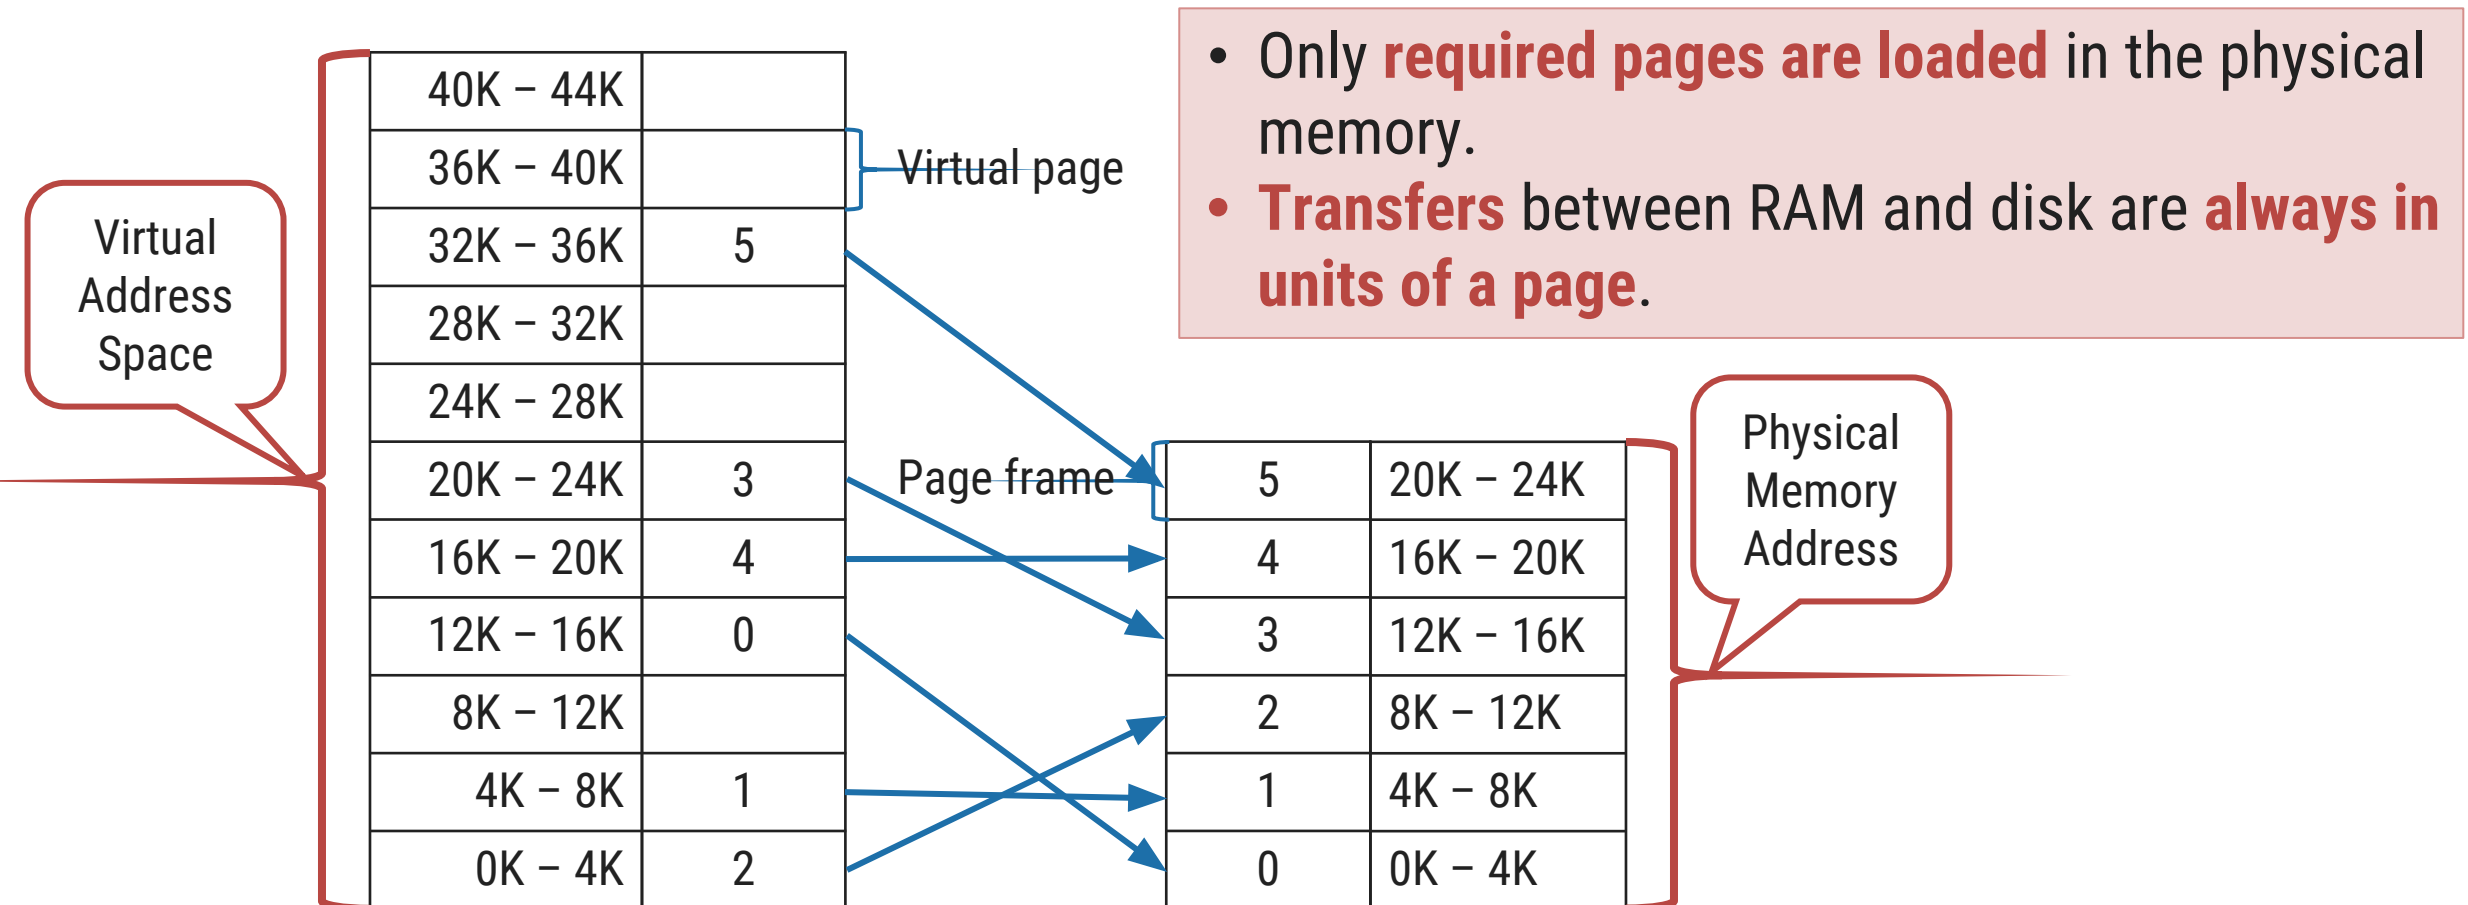
Conversion of virtual address to physical address

- When virtual memory is used, the **virtual address is presented to an MMU** (Memory Management Unit) that **maps the virtual addresses onto the physical memory addresses**.
- We have a computer generated 16-bit addresses, from 0 up to 44K. These are the virtual addresses.

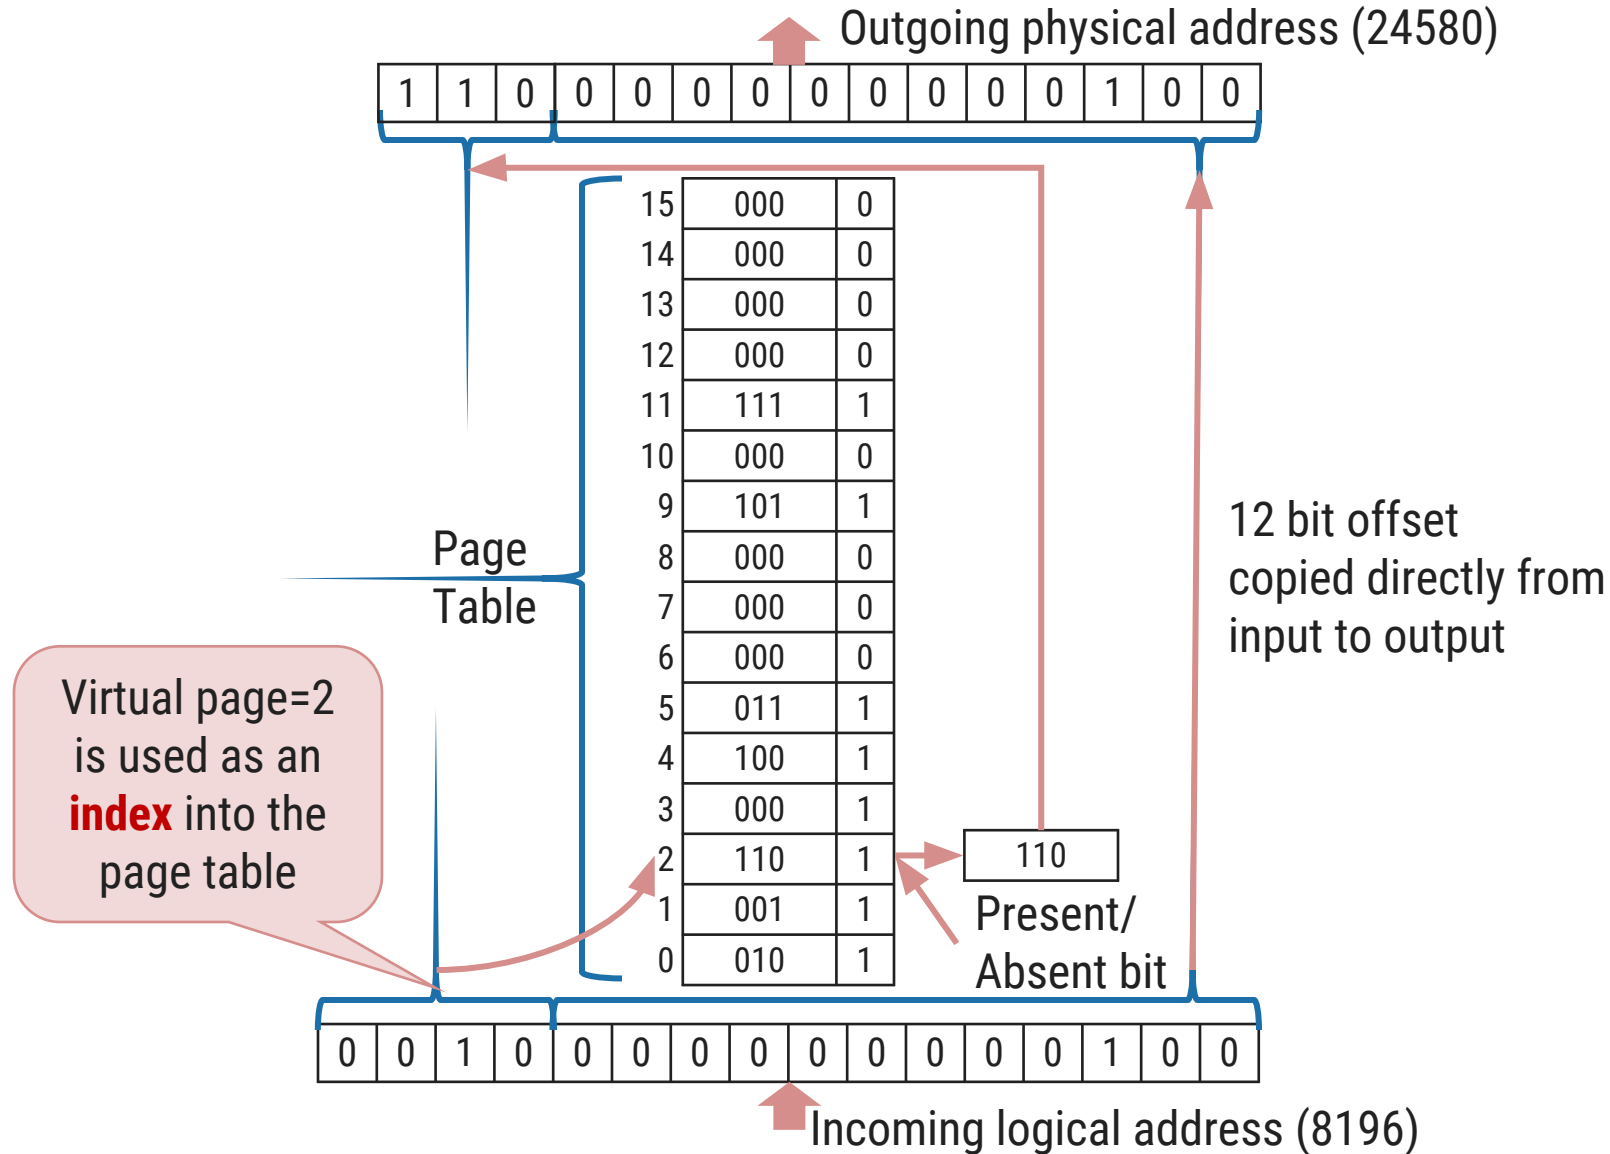


Conversion of virtual address to physical address

- However, only **24KB of physical memory is available**, so although **44KB programs can be written, they cannot be loaded in to memory in their entirety and run.**
- A complete copy of a program's core image, up to 44 KB, must be present on the disk.



Internal operation of the MMU




Internal operation of the MMU

- The **virtual address is split into a virtual page number** (high order bits) and **an offset** (low-order bits).
- With a 16-bit address and a 4KB page size, the **upper 4 bits could specify one of the 11 virtual pages** and the **lower 12 bits would then specify the byte offset** (0 to 4095) within the selected page.
- The **virtual page number is used as an index** into the Page table.
- If the **present/absent bit is 0**, it is **page-fault**; a trap to the **operating system is caused to bring required page into main memory**.
- If the **present/absent bit is 1**, **required page is there with main memory** and **page frame number found in the page table is copied to the higher order bit of the output register along with the offset**.
- Together **Page frame number and offset creates physical address**.
- **Physical Address = Page frame Number + offset of virtual address**.

Page table

- ❑ Page table is a **data structure which translates virtual address into equivalent physical address.**
- ❑ The virtual page number is used as an index into the Page table to find the entry for that virtual page and from the Page table physical page frame number is found.
- ❑ Thus the **purpose of page table is to map virtual pages onto page frames.**
- ❑ The page table of the process is **held in the kernel space.**

Page table



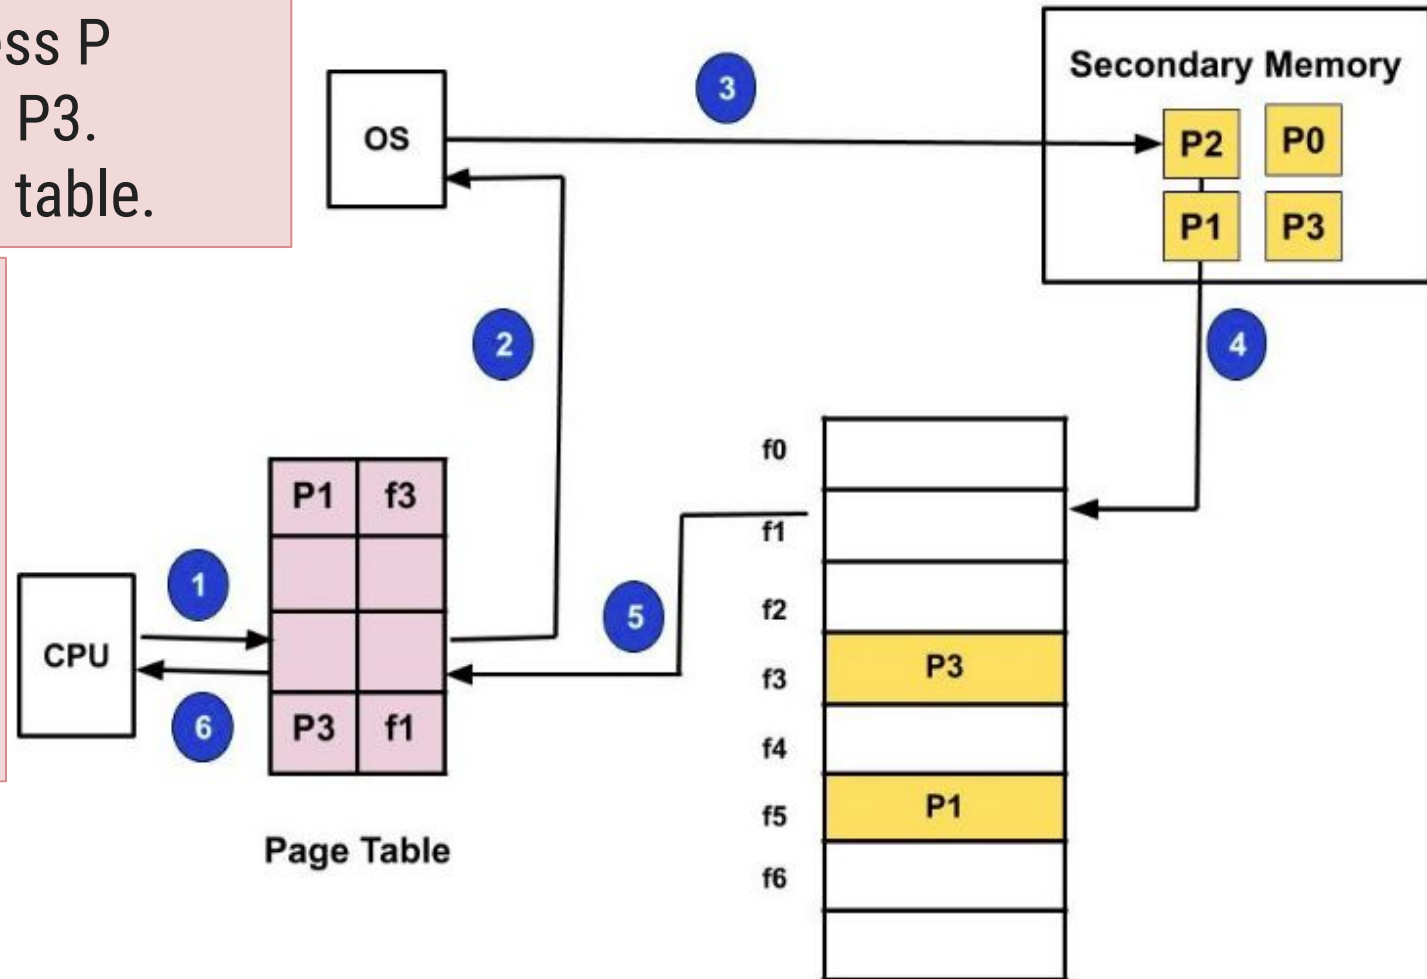
- ❑ **Page frame Number:** It gives the **frame number in which the current page you are looking for is present**.
- ❑ **Present/Absent bit:** Present or absent bit **says whether a particular page you are looking for is present or absent**. If it is not present, that is called Page Fault. It is set to 0 if the corresponding page is not in memory. Sometimes this bit is also known as valid/invalid bits.
- ❑ **Protection bits:** Protection bit **says that what kind of protection you want on that page**. In the simplest form, 0 for read/write and 1 for read only.
- ❑ **Modified bit:** Modified bit **says whether the page has been modified or not**. If the page in memory has been modified, it must be written back to disk. This bit is also called as dirty bit as it reflects the page's state.
- ❑ **Referenced bit:** A references bit is **set whenever a page is referenced, either for reading or writing**. Its value helps operating system in page replacement algorithm.
- ❑ **Cashing Disabled bit:** This feature is important for pages that **maps onto device registers rather than memory**. With this bit caching can be turned off.

Definitions (Demand paging)

□ **Demand paging:** Demand paging is a **technique used in virtual memory** systems where the **pages are brought in the main memory only when required or demanded by the CPU.**

- Suppose we have to execute a process P having four pages as P0, P1, P2, and P3.
- We have page P1 and P3 in the page table.

- Now, if the CPU wants to access page P2 of a process P
 - search the page in the page table
 - P2 is not available in page table so page fault
 - OS bring P2 from HDD to RAM
 - OS will update page table
 - Process P will be executed.



Demand paging

- Now, if the CPU wants to access page P2 of a process P, first it will search the page in the page table.
- As the page table does not contain this page so it will be a trap or page fault. As soon as the trap is generated and context switching happens and the control goes to the operating system.
- The OS system will put the process in a waiting/ blocked state. The OS system will now search that page in the backing store or secondary memory.
- The OS will then read the page from the backing store and load it to the main memory.
- Next, the OS system will update the page table entry accordingly.
- Finally, the control is taken back from the OS and the execution of the process is resumed.
- Hence whenever a page fault occurs these steps are followed by the operating system and the required page is brought into memory.

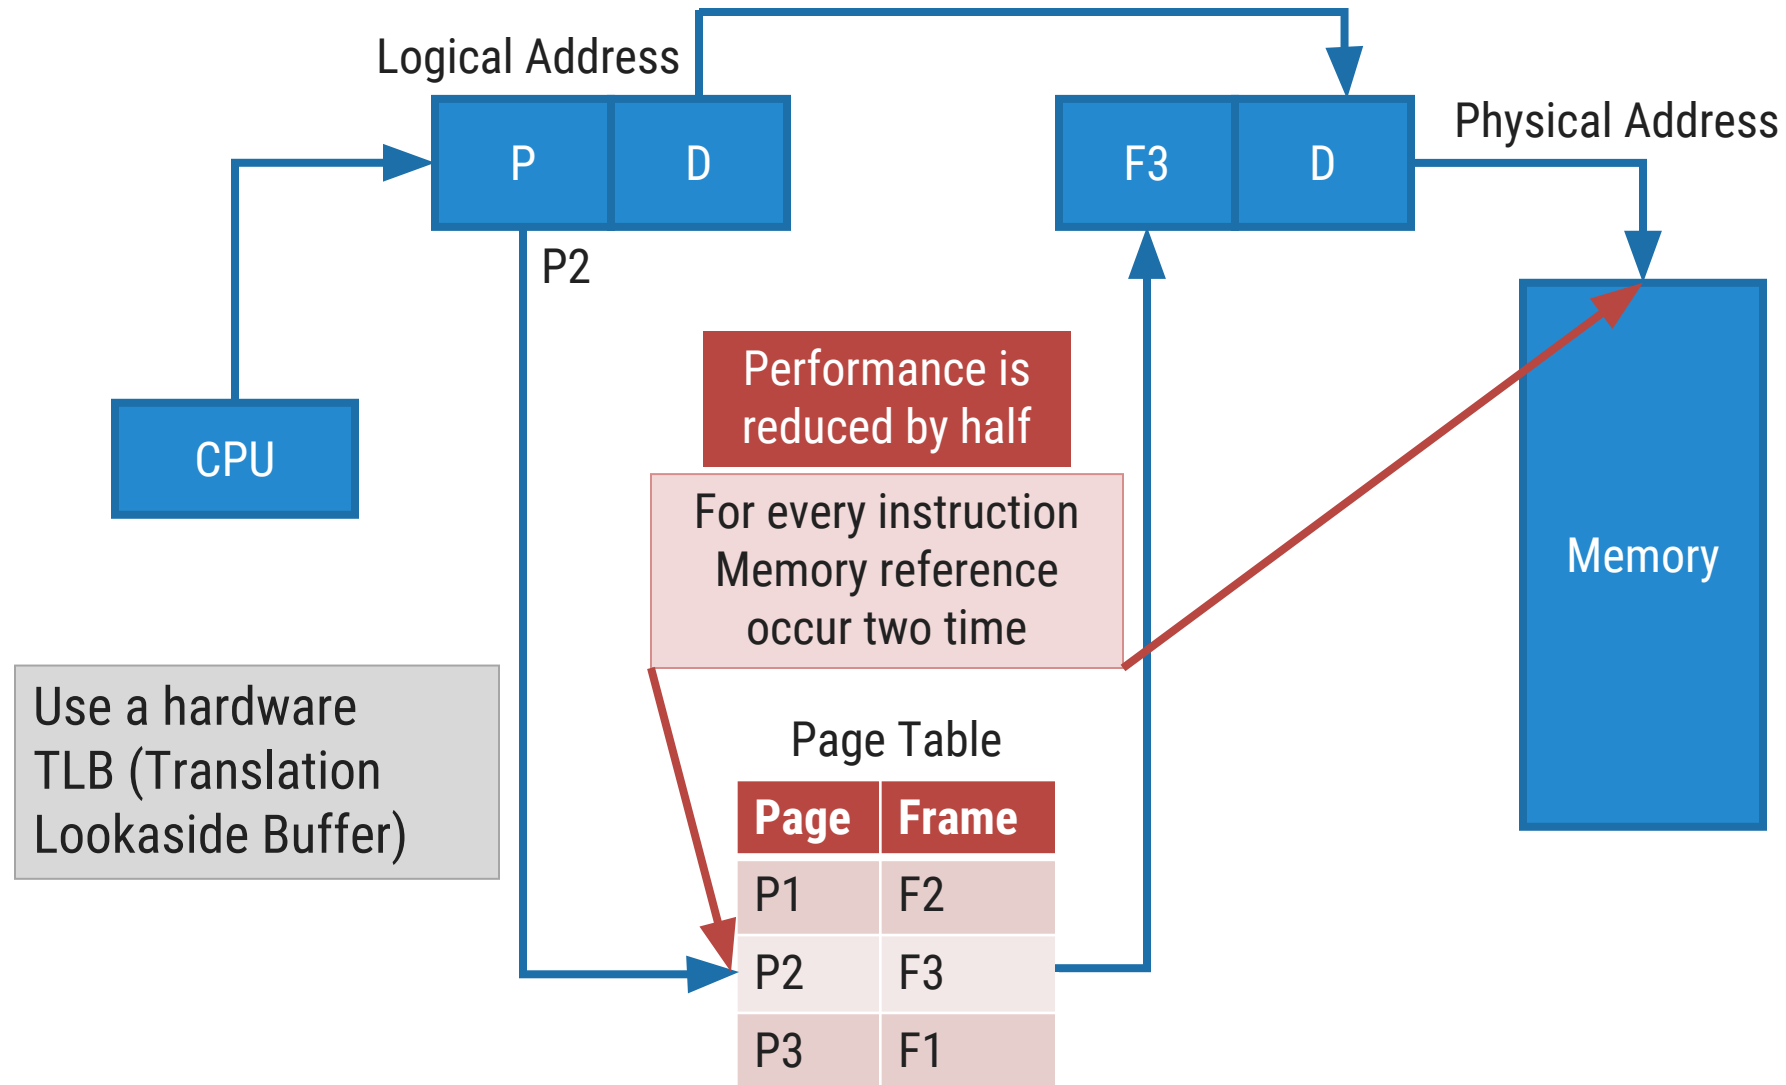
Definitions

- ❑ **Pre-paging**: Many paging systems try to keep track of each process's working set and make sure that it is in memory before letting the process run. **Loading pages before allowing processes to run** is called pre-paging.
- ❑ **Working Set**: The **set of pages that a process is currently using** is known as working set.
- ❑ **Thrashing**: Thrashing is a condition or a situation **when the system is spending a major portion of its time in servicing the page faults**, but the actual processing done is very negligible.

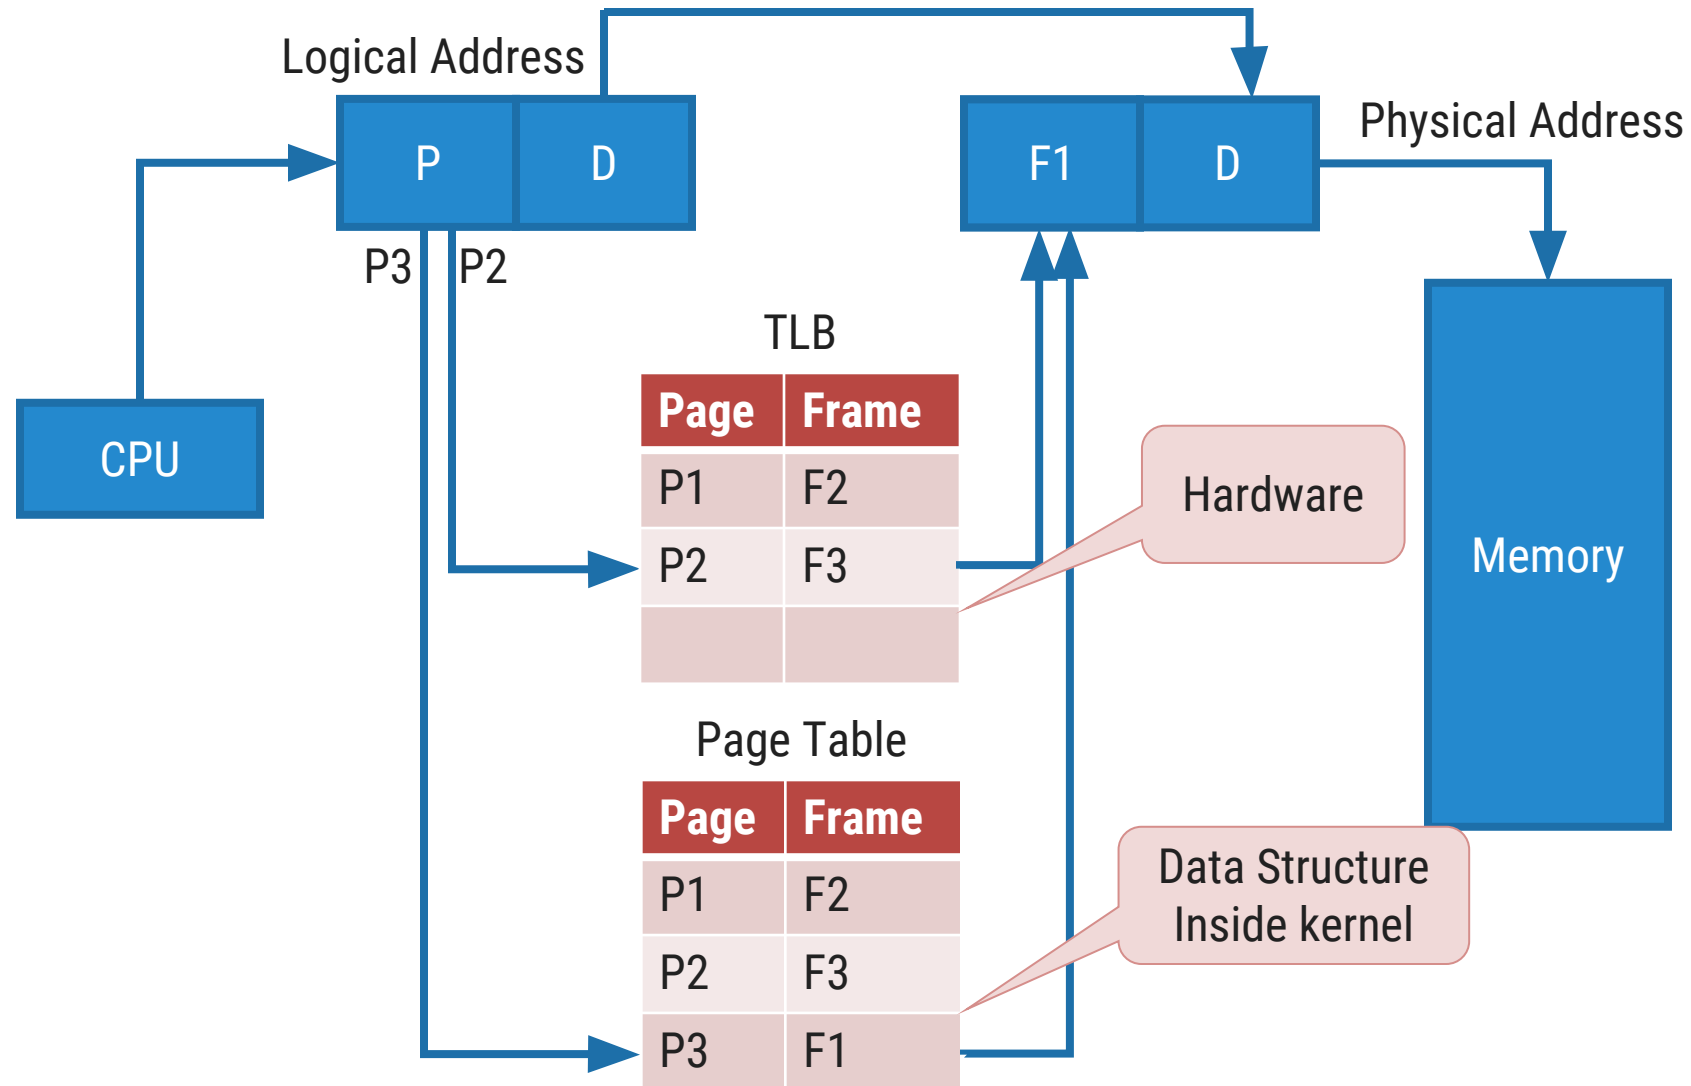
Issues in Paging

- In any paging system, two major issues must be faced:
 1. The **mapping from virtual address to physical address must be fast.**
 2. If the **virtual address space is large, the page table will be large.**

Mapping from virtual address to physical address must be fast



Mapping from virtual address to physical address must be fast



Mapping from virtual address to physical address using TLB

□ Steps in TLB hit:

- CPU generates virtual address.
- It is checked in TLB (present).
- Corresponding frame number is retrieved, which now tells where in the main memory page lies.

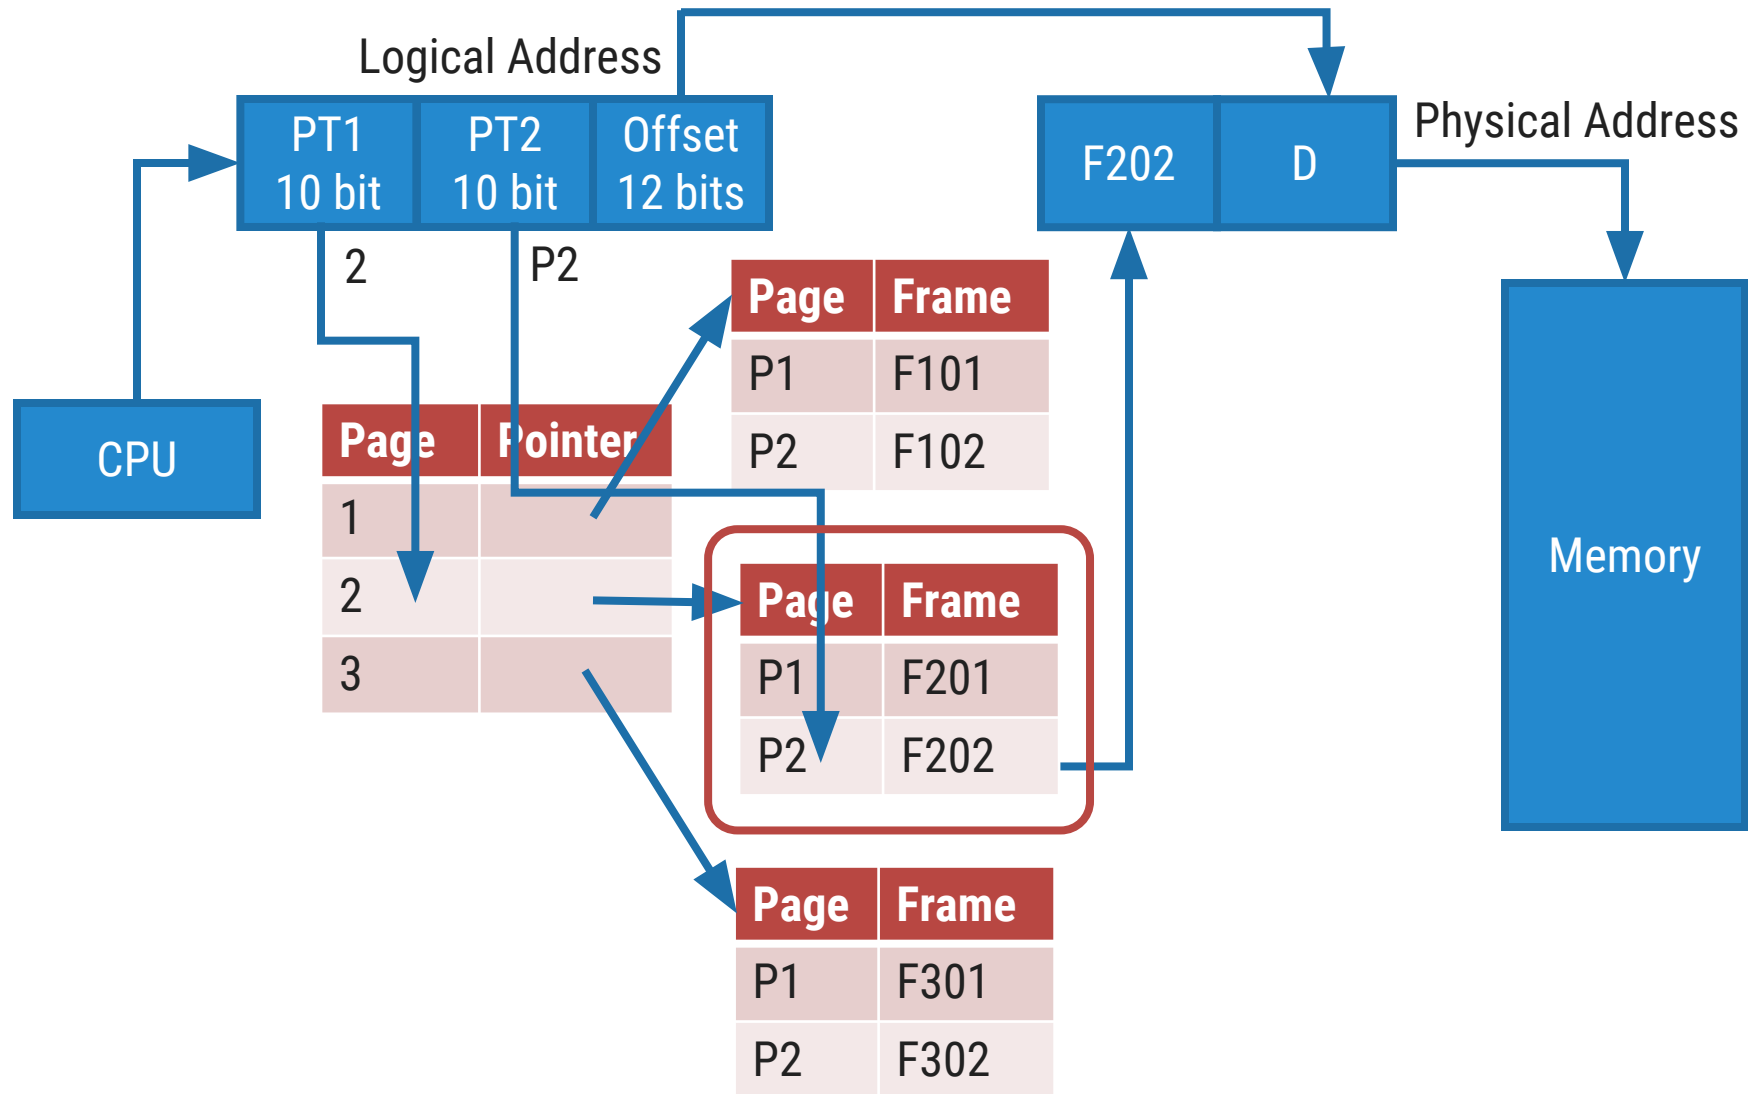
□ Steps in Page miss:

- CPU generates virtual address.
- It is checked in TLB (not present).
- Now the page number is matched to page table residing in main memory.
- Corresponding frame number is retrieved, which now tells where in the main memory page lies.
- The TLB is updated with new Page Table Entry (if space is not there, one of the replacement technique comes into picture i.e either FIFO, LRU or MFU etc).

Virtual address space is large, the page table will be large


- Two different ways to deal with large page table problems:
 1. Multilevel Page Table
 2. Inverted Page Table

Multilevel Page Table



Inverted Page Table

Current process ID - 245



Process IDs do not match.
So follow chaining pointer

Frame No	Page No	Process ID	Pointer
1	2	211	
2			
3			
4	2	245	
5			
6			

Process IDs match.
So frame no added to
physical address.

Inverted Page Table

- Each entry in the page table contains the following fields.
 - Page number – It specifies the page number range of the logical address.
 - Process id – An inverted page table contains the address space information of all the processes in execution.
 - Since two different processes can have similar set of virtual addresses, it becomes necessary in Inverted Page Table to store a processID of each process to identify its address space uniquely.
 - This is done by using the combination of PID and Page Number. So this Process Id acts as an address space identifier and ensures that a virtual page for a particular process is mapped correctly to the corresponding physical frame.
 - Control bits – These bits are used to store extra paging-related information. These include the valid bit, dirty bit, reference bits, protection and locking information bits.
 - Chained pointer – It may be possible sometime that two or more processes share a part of main memory.
 - In this case, two or more logical pages map to same Page Table Entry then a chaining pointer is used to map the details of these logical pages to the root page table.



Page Replacement Algorithms

Section - 7



Page replacement algorithms

□ Following are different types of page replacement algorithms

1. Optimal Page Replacement Algorithm
2. FIFO Page Replacement Algorithm
3. The Second Chance Page Replacement Algorithm
4. The Clock Page Replacement Algorithm
5. LRU (Least Recently Used) Page Replacement Algorithm
6. NRU (Not Recently Used)

Optimal Page Replacement Algorithm

- The moment a page fault occurs, some set of pages will be in the memory.
- One of these pages will be referenced on the very next instruction.
- Other pages may not be referenced until 10, 100, or perhaps 1000 instructions later.
- Each **page can be labeled with the number of instructions that will be executed before that page is first referenced.**
- The optimal page algorithm simply says that the **page with the highest label should be removed.**
- The only problem with this algorithm is that it is **unrealizable.**
- At the time of the page fault, the **operating system has no way of knowing when each of the pages will be referenced next.**

Optimal Page Replacement Algorithm

□ Page Reference String:

- 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 0, 2, 0, 1, 7, 0, 1
- Three frames

Page Request	7	0	1	2	0	3	0	4	2	3	0	3	2	0	2	0	1	7	0	1
Frame – 1	7	7	7	2	2	2	2	2	2	2	2	2	2	2	2	2	1	1	1	1
Frame – 2		0	0	0	0	0	0	4	4	4	0	0	0	0	0	0	0	0	0	0
Frame – 3			1	1	1	3	3	3	3	3	3	3	3	3	3	3	3	7	7	7
Page Faults (9)	F	F	F	F		F		F			F						F	F		

FIFO Page Replacement Algorithm

- ❑ The first in first out page replacement algorithm is the simplest page replacement algorithm.
- ❑ The operating system **maintains a list of all pages currently in memory, with the most recently arrived page at the tail and least recent at the head.**
- ❑ On a page fault, the **page at head is removed and the new page is added to the tail.**
- ❑ When a page replacement is **required the oldest page in memory needs to be replaced.**
- ❑ The performance of the FIFO algorithm is **not always good because it may happen that the page which is the oldest is frequently referred by OS.**
- ❑ Hence removing the **oldest page may create page fault again.**

FIFO Page Replacement Algorithm


□ Page Reference String:

- 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1
- Three frames

Page Request	7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
Frame – 1	7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
Frame – 2		0	0	0	0	3	3	3	2	2	2	2	2	1	1	1	1	1	0	0
Frame – 3			1	1	1	1	0	0	0	3	3	3	3	3	2	2	2	2	2	1
Page Faults (15)	F	F	F	F		F	F	F	F	F	F			F	F			F	F	F


Second Chance Page Replacement Algorithm

- It is **modified form of the FIFO page replacement algorithm**.
- It looks at the front of the queue as FIFO does, but instead of immediately paging out that page, it **checks to see if its referenced bit is set**.
 - If it is not set (zero), the page is swapped out.
 - Otherwise, the referenced bit is cleared, the **page is inserted at the back of the queue** (as if it were a new page) and this process is repeated.



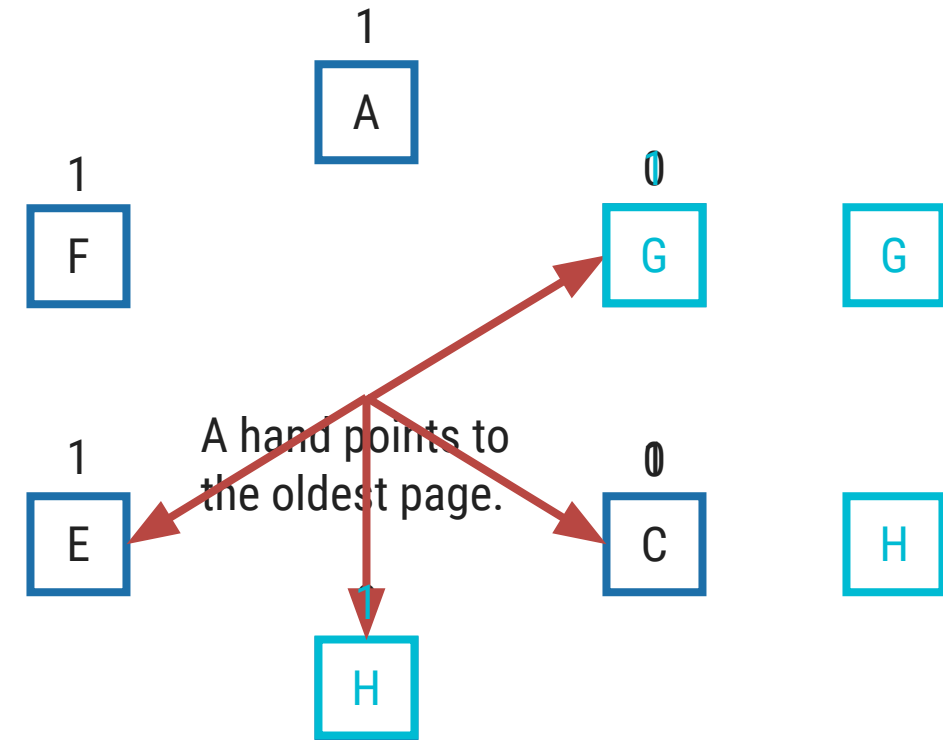
Second Chance Page Replacement Algorithm

- If all the pages have their referenced bit set, on the second encounter of the first page in the list, that page will be swapped out, as it now has its referenced bit cleared.
- If all the pages have their reference bit cleared, then second chance algorithm degenerates into pure FIFO.



Clock Page Replacement Algorithm

- In second chance algorithm **pages are constantly moving around on its list**. So it is **not working efficiently**.
- A better approach is to **keep all the page frames on a circular list in the form of a clock**.
- When a page fault occurs, the **page being pointed to by the hand is inspected**.
 - If its **R bit is 0**, the page is evicted, the **new page is inserted into the clock in its place**, and the **hand is advanced one position**.
 - If **R is 1**, it is **cleared** and the **hand is advanced to the next page**.



LRU (Least Recently Used) Page Replacement Algorithm

- ❑ A good approximation to the optimal algorithm is based on the observation that pages that have been heavily used in last few instructions will probably be heavily used again in next few instructions.
- ❑ When page fault occurs, **throw out the page that has been used for the longest time**. This strategy is called LRU (Least Recently Used) paging.
- ❑ To fully implement LRU, it is necessary to maintain a linked list of all pages in memory, with the **most recently used page at the front and the least recently used page at the rear**.
- ❑ The list must be updated on every memory reference.
- ❑ Finding a page in the list, deleting it, and then moving it to the front is a very time consuming operations.

LRU (Least Recently Used) Page Replacement Algorithm

□ Page Reference String:

- 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1
- Three frames

Page Request	7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
Frame – 1	7	7	7	2	2	2	2	4	4	4	0	0	0	1	1	1	1	1	1	1
Frame – 2		0	0	0	0	0	0	0	0	3	3	3	3	3	3	0	0	0	0	0
Frame – 3			1	1	1	3	3	3	2	2	2	2	2	2	2	2	2	7	7	7
Page Faults (12)	F	F	F	F		F		F	F	F	F			F		F		F		

NRU (Not Recently Used) Page Replacement Algorithm

- ❑ NRU makes approximation to replace the page based **on R (referenced) and M (modified) bits**.
- ❑ When a process is started up, **both page bits for all pages are set to 0 by operating system**.
- ❑ Periodically, the **R bit is cleared**, to distinguish pages that have not been referenced recently from those that have been.
- ❑ When page fault occurs, the operating system inspects all the pages and **divide them into 4 categories based on current values of their R and M bits**
 - ❑ Case 0 : not referenced, not modified
 - ❑ Case 1 : not referenced, modified
 - ❑ Case 2 : referenced, not modified
 - ❑ Case 3 : referenced, modified
- ❑ The NRU (Not Recently Used) algorithm removes a page at random from the lowest numbered nonempty class.

NRU (Not Recently Used) Page Replacement Algorithm

- For example if,
 - Page-0 is of class-2 (referenced, not modified)
 - Page-1 is of class-1 (not referenced, modified)
 - Page-2 is of class-0 (not referenced, not modified)
 - Page-3 is of class-3 (referenced, modified)
- So lowest class **page-2 needs to be replaced by NRU.**

Belady's Anomaly (in FIFO Page Replacement Algorithm)

□ Page Reference String: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Page Faults of 4 Frame > Page Faults of 3 Frame

Three
Frames

Page Request	1	2	3	4	1	2	5	1	2	3	4	5
Frame – 1	1	1	1	4	4	4	5	5	5	5	5	5
Frame – 2		2	2	2	1	1	1	1	1	3	3	3
Frame – 3			3	3	3	2	2	2	2	2	4	4
Page Faults (9)	F	F	F	F	F	F	F			F	F	

Four
Frames

Page Request	1	2	3	4	1	2	5	1	2	3	4	5
Frame – 1	1	1	1	1	1	1	5	5	5	5	4	4
Frame – 2		2	2	2	2	2	2	1	1	1	1	5
Frame – 3			3	3	3	3	3	3	2	2	2	2
Frame – 4				4	4	4	4	4	4	3	3	3
Page Faults (10)	F	F	F	F			F	F	F	F	F	F

Belady's Anomaly (in FIFO Page Replacement Algorithm)

- ❑ Belady's anomaly is the **phenomenon in which increasing the number of page frames results in an increase in the number of page faults for certain memory access patterns.**
- ❑ This phenomenon is commonly **experienced when using the First-In First-Out (FIFO) page replacement algorithm.**
- ❑ In FIFO, the page fault may or may not increase as the page frames increase, but in Optimal and stack-based algorithms like LRU, as the page frames increase the page fault decreases.

Sum

- A computer has four page frames. The time of loading, time of last access and the R and M bit for each page given below. Which page FIFO, LRU and NRU will replace.

Page	Loaded	Last used	R	M
0	126	280	1	0
1	230	265	0	1
2	140	270	0	0
3	110	285	1	1

- **FIFO**:- Page which is **arrived first needs to be removed first**, so **page-3** needs to replace.

- **LRU**:- When page fault occurs, throw out the **page that has been used for the longest time**.

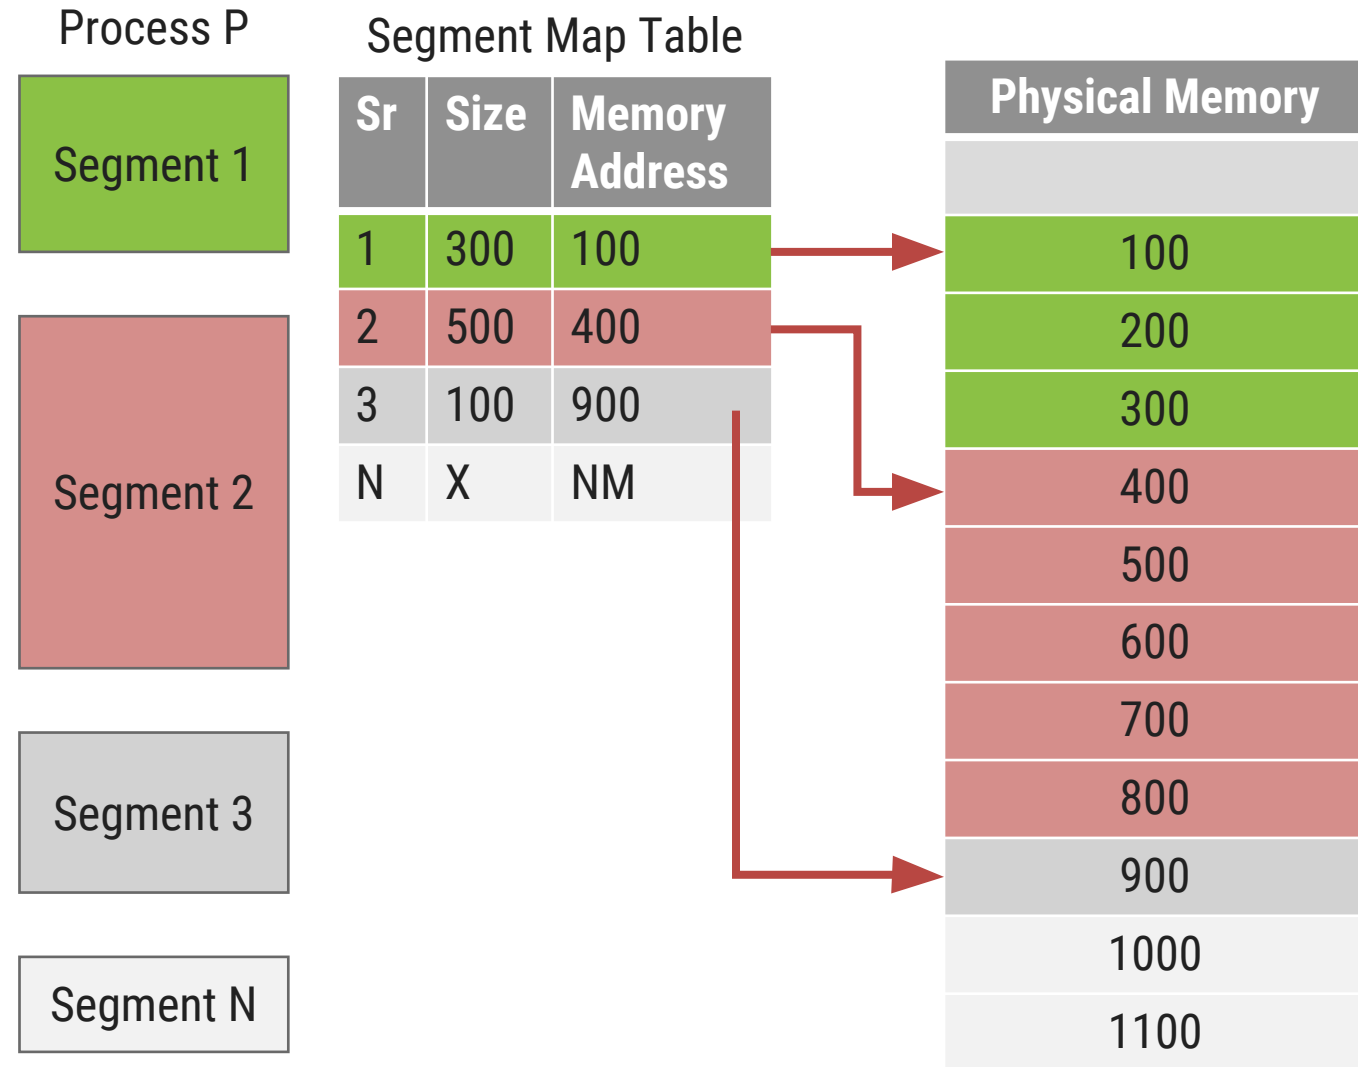
- **Page-1** is not used for the long time from all four, so LRU suggest replacing page-1.

- **NRU**:-

- Page-0 is of class-2 (referenced, not modified)
- Page-1 is of class-1 (not referenced, modified)
- Page-2 is of class-0 (not referenced, not modified)
- Page-3 is of class-3 (referenced, modified)
- So lowest class **page-2** needs to be replaced by NRU

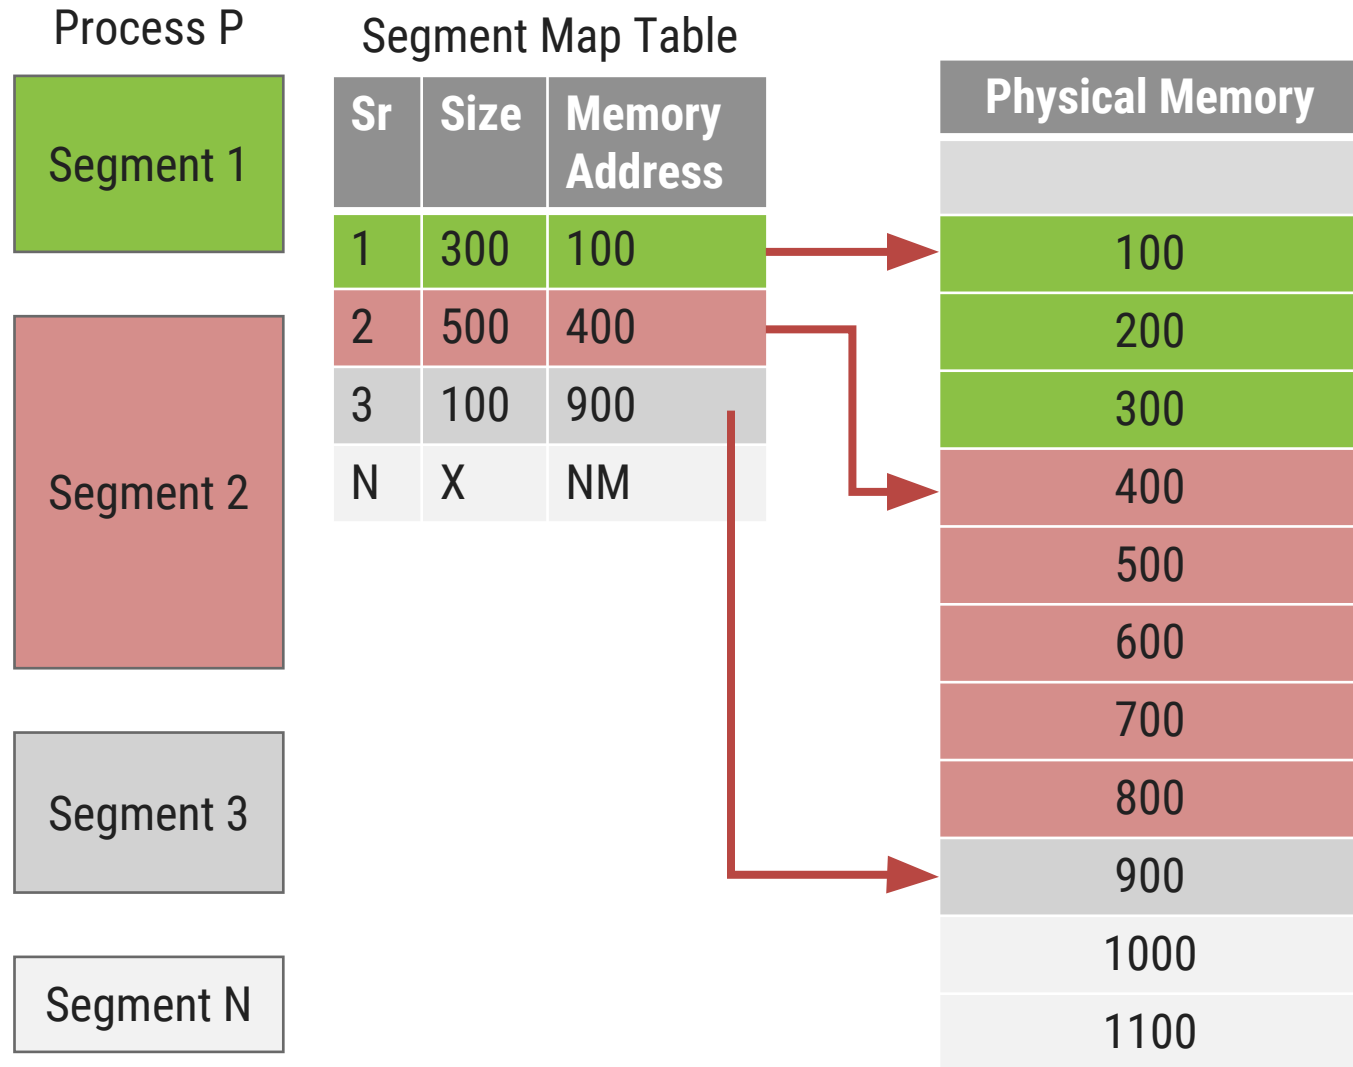
Segmentation

- ❑ Segmentation is a **memory management technique in which each job is divided into several segments of different sizes**, one for each module that contains pieces that perform related functions.
- ❑ **Each segment is actually a different logical address space** of the program.
- ❑ **When a process is to be executed, its corresponding segmentation are loaded into non-contiguous memory** though every segment is loaded into a contiguous block of available memory.
- ❑ Segmentation memory management works very similar to paging but here **segments are of variable-length where as in paging pages are of fixed size**.



Segmentation

- A program segment contains the program's main function, utility functions, data structures, and so on.
- The **operating system maintains a segment map table for every process.**
- **Segment map table contains list of free memory blocks along with segment numbers, their size and corresponding memory locations in main memory.**
- For each segment, the **table stores the starting address of the segment and the length of the segment.**
- A reference to a memory location includes a value that identifies a segment and an offset.



Paging VS Segmentation

Paging	Segmentation
Paging was invented to get large address space without having to buy more physical memory.	Segmentation was invented to allow programs and data to be broken up into logically independent address space and to add sharing and protection.
The programmer does not aware that paging is used.	The programmer is aware that segmentation is used.
Procedure and data cannot be distinguished and protected separately.	Procedure and data can be distinguished and protected separately.
Change in data or procedure requires compiling entire program.	Change in data or procedure requires compiling only affected segment not entire program.
Sharing of different procedures not available.	Sharing of different procedures available.

Questions

1. Explain multiprogramming with fixed partition.
2. Explain link list method for dynamic memory management.
3. How free space can be managed by OS.
4. Explain Swapping and Fragmentation in detail.
5. What is Paging? Explain paging mechanism in MMU with example.
6. Explain TLB and Virtual Memory.
7. Discuss demand paging.
8. Define following Terms: Thrashing
9. List different Page Replacement Algorithms? Discuss it in terms of page faults.
10. What is Belady's anomaly? Explain with suitable example.
11. Consider (70120304230321201701) page reference string: How many page fault would occur for following page replacement algorithm. Consider 3 frames and 4 frames.
 1. FIFO
 2. LRU
 3. Optimal

Questions asked in GTU

12. Given six Partition of 300KB, 600KB, 350KB, 200KB, 750KB and 125KB(in order), how would the first-fit, best-fit and worst-fit algorithms places processes of size 115 KB, 500KB, 358KB, 200KB and 375KB(in order)? Which algorithm is efficient for the use of memory?
13. Calculate the page fault rates for below reference string in case of FIFO and Optimal page replacement algorithm. Assume the memory size is 4 page frames and all frames are initially empty. 0,2,1,6,4,0,1,0,3,1,2,1
14. Consider the following reference string. Calculate the page fault rates for FIFO and OPTIMAL page replacement algorithm. Assume the memory size is 4 page frame.
1,2,3,4,5,3,4,1,6,7,8,7,8,9,7,8,9,5,4,5,4,2
15. Consider the page reference string: 1,2,3,4,5,3,4,1,6,7,8,7,8,9,7,8,9,5,4,5,4,2 With four Frames How many page faults would occur for the FIFO, Optimal page replacement algorithms? which algorithm is efficient? (assume all frames are initially empty)